

# 432 Class 03

Thomas E. Love, Ph.D.

2026-01-20

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

1

## Today's Agenda

- A First Example: Space Shuttle O-Rings
- Predicting a Binary outcome using a single predictor
  - using a linear probability model
  - using logistic regression and `glm`

See Chapters 19-20 in our [Course Notes](#) for more on logistic regression and related models.

# Today's R Setup

```
1 knitr::opts_chunk$set(comment = NA)
2
3 library(janitor)
4 library(naniar)
5
6 library(broom)
7 library(caret) # for confusion matrix
8 library(faraway) # data source
9 library(gt)
10 library(patchwork)
11
12 library(easystats)
13 library(tidyverse)
14
15 theme_set(theme_bw())
```

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

3

## Challenger Space Shuttle Data

The US space shuttle Challenger exploded on 1986-01-28. An investigation ensued into the reliability of the shuttle's propulsion system. The explosion was eventually traced to the failure of one of the three field joints on one of the two solid booster rockets. Each of these six field joints includes two O-rings which can fail.

- The discussion among engineers and managers raised concern that the probability of failure of the O-rings depended on the temperature at launch, which was forecast to be 31 degrees F.
- There are strong engineering reasons based on the composition of O-rings to support the judgment that failure probability may rise monotonically as temperature drops.

We have data on 23 space shuttle flights that preceded *Challenger* on primary O-ring erosion and/or blowby and on the temperature in degrees Fahrenheit. No previous liftoff temperature was under 53 degrees F.

# The “O-rings” data

- `damage` = number of damage incidents out of 6 possible
- we set `burst` = 1 if `damage` > 0

```
1 orings1 <- faraway::orings |> tibble() |>
2   mutate(burst = case_when( damage > 0 ~ 1, TRUE ~ 0))
3
4 orings1 |> summary()
```

	temp	damage	burst
Min.	:53.00	Min. :0.0000	Min. :0.0000
1st Qu.	:67.00	1st Qu.:0.0000	1st Qu.:0.0000
Median	:70.00	Median :0.0000	Median :0.0000
Mean	:69.57	Mean :0.4783	Mean :0.3043
3rd Qu.	:75.00	3rd Qu.:1.0000	3rd Qu.:1.0000
Max.	:81.00	Max. :5.0000	Max. :1.0000

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

5


## Association of `burst` and `temp`

```
1 ggplot(orings1, aes(x = factor(burst), y = temp)) +
2   geom_violin() +
3   geom_boxplot(aes(fill = factor(burst)), width = 0.3) +
4   stat_summary(geom = "point", fun = mean, col = "white", size = 2.5) +
5   guides(fill = "none") +
6   labs(title = "Are bursts more common at low temperatures?",
7        subtitle = "23 prior space shuttle launches",
8        x = "Was there a burst? (1 = yes, 0 = no)",
9        y = "Launch Temp (F)")
```

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

6

# Association of `burst` and `temp`



432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

7

## Predict Prob(burst) using temperature?

We want to treat the binary variable `burst` as the outcome, and `temp` as the predictor.


- We'll jitter the points vertically so that they don't overlap completely if we have two launches with the same temperature.

```
1 ggplot(orings1, aes(x = temp, y = burst)) +  
2   geom_jitter(col = "navy", size = 3, width = 0, height = 0.1) +  
3   labs(title = "Are bursts more common at low temperatures?",  
4         subtitle = "23 prior space shuttle launches",  
5         y = "Was there a burst? (1 = yes, 0 = no)",  
6         x = "Launch Temp (F)")
```

# Predict Prob(burst) using temperature?

Are bursts more common at low temperatures?

23 prior space shuttle launches



# A Linear Probability Model, fit with `lm()`

# Linear model to predict Prob(burst)?

```
1 fit1 <- lm(burst ~ temp, data = orings1)
2
3 tidy(fit1, conf.int = T) |> gt() |>
4   fmt_number(decimals = 3) |> tab_options(table.font.size = 20)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	2.905	0.842	3.450	0.002	1.154	4.656
temp	-0.037	0.012	-3.103	0.005	-0.062	-0.012

- This is a linear probability model.


$$\widehat{\text{burst}} = 2.905 - 0.037(\text{temp})$$

## Plot linear probability model?

```
1 ggplot(orings1, aes(x = temp, y = burst)) +
2   geom_jitter(col = "navy", size = 3, width = 0, height = 0.1) +
3   geom_smooth(method = "lm", se = F, col = "red",
4               formula = y ~ x) +
5   labs(title = "Bursts more common at lower temperatures",
6        subtitle = "23 prior space shuttle launches",
7        y = "Was there a burst? (1 = yes, 0 = no)",
8        x = "Launch Temp (F)")
```

- It would help if we could see the individual launches...

# Plot linear probability model?



432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

13

## Making Predictions with `fit1`

```
1 fit1$coefficients
```

```
(Intercept)      temp
2.90476190 -0.03738095
```

- What does `fit1` predict for the probability of a burst if the temperature at launch is 70 degrees F?

```
1 predict(fit1, newdata = tibble(temp = 70))
```

```
1
0.2880952
```

- What if the temperature was actually 60 degrees F?

# Making Predictions with `fit1`

Let's use our linear probability model `fit1` to predict the probability of a burst at some other temperatures...

```
1 newtemps <- tibble(temp = c(80, 70, 60, 50, 31))
2
3 augment(fit1, newdata = newtemps)


# A tibble: 5 × 2
  temp .fitted
  <dbl>   <dbl>
1     80 -0.0857
2     70  0.288
3     60  0.662
4     50  1.04
5     31  1.75
```

- Uh, oh.


## Checking model `fit1` (1/2)

```
1 check_model(fit1, detrend = FALSE, check = c("pp_check", "qq"))
```

Posterior Predictive Check  
Model-predicted lines should resemble observed data line




Normality of Residuals  
Dots should fall along the line




# Checking model `fit1` (2/2)

```
1 check_model(fit1, detrend = FALSE, check = c("linearity", "homogeneity"))
```

Linearity  
Reference line should be flat and horizontal



Homogeneity of Variance  
Reference line should be flat and horizontal



## Models to predict a Binary Outcome

Our outcome takes on two values (zero or one) and we then model the probability of a “one” response given a linear function of predictors.

Idea 1: Use a *linear probability model*

- Main problem: predicted probabilities that are less than 0 and/or greater than 1
- Also, how can we assume Normally distributed residuals when outcomes are 1 or 0?

# Models to predict a Binary Outcome

Idea 2: Build a *non-linear* regression approach

- Most common approach: logistic regression, part of the class of *generalized* linear models

# A Logistic Regression Model, fit with `glm()`

# The Logit Link and Logistic Function

The function we use in logistic regression is called the **logit link**.


The inverse of the logit function is called the **logistic function**.  
 $\text{logit}(\pi) = \log\left(\frac{\pi}{1-\pi}\right) = \beta_0 + \beta_1 x$

If  $\text{logit}(\pi) = \eta$ , then .

$$\pi = \frac{e^{\eta}}{1+e^{\eta}}$$

- The logistic function takes any value in the real numbers and returns a value between 0 and 1.  $x$

## The Logistic Function



# The logit or log odds

We usually focus on the **logit** in statistical work, which is the inverse of the logistic function.


- If we have a probability , then .  
 $\pi < \text{logit}(\pi) < 0$
- If our probability , then .  
 $\pi > \text{logit}(\pi) > 0$
- Finally, if , then .  
 $\pi = \text{logit}(\pi) = 0$

## Why is this helpful?

- $\log(\text{odds}(Y = 1))$  or  $\text{logit}(Y = 1)$  covers all real numbers.
- $\text{Prob}(Y = 1)$  is restricted to  $[0, 1]$ .

# Predicting Pr(event) or Pr(no event)

- Can we flip the story?



# Back to predicting Prob(burst)

We'll use the `glm` function in R, specifying a logistic regression model.

- Instead of predicting , we're predicting or .

$$Pr(\text{burst}) \quad \log(\text{odds}(\text{burst}))$$

## fit2 for Prob(burst)

```
1 fit2 <- glm(burst ~ temp, data = orings1,
2               family = binomial(link = "logit"))
3
4 tidy(fit2, conf.int = TRUE) |> gt() |>
5   fmt_number(decimals = 3) |> tab_options(table.font.size = 24)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	15.043	7.379	2.039	0.041	3.331	34.342
temp	-0.232	0.108	-2.145	0.032	-0.515	-0.061

$$\log \left[ \frac{\widehat{P(\text{burst} = 1)}}{1 - \widehat{P(\text{burst} = 1)}} \right] = 15.0$$

# Understanding **fit2**'s predictions

- For a temperature of 70 F at launch, what is our prediction?
  - $\log(\text{odds}(\text{burst})) = 15.043 - 0.232(70) = -1.197$
  - $\text{odds}(\text{burst}) = \exp(-1.197) = 0.302$
  - so, we can estimate the probability by

$$Pr(\text{burst}) = \frac{0.302}{(0.302 + 1)} = 0.25$$

## Prediction from **fit2** for temp = 60

What is the predicted probability of a burst if the temperature is 60 degrees?

- $\log(\text{odds}(\text{burst})) = 15.043 - 0.232(60) = 1.123$
- $\text{odds}(\text{burst}) = \exp(1.123) = 3.074$
- $Pr(\text{burst}) = 3.074 / (3.074 + 1) = 0.755$

# Using predict(fit2)

What is the predicted probability of a burst?

```
1 temps <- tibble(temp = c(40,50,60,70,80))
2
3 predict(fit2, newdata = temps, type = "link") # est. log odds of burst
```

	1	2	3	4	5
5.756392	3.434764	1.113137	-1.208490	-3.530118	

```
1 predict(fit2, newdata = temps, type = "response") # fitted Pr(burst)
```

	1	2	3	4	5
0.99684747	0.96877352	0.75271348	0.22996826	0.02846733	

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

29

## Will augment do this, as well?

Yes, and it will retain many more decimal places in intermediate calculations...

```
1 temps <- tibble(temp = c(60,70))
2
3 augment(fit2, newdata = temps, type.predict = "link")
```

	# A tibble: 2 × 2	temp .fitted
1	<dbl>	<dbl>
2	60	1.11
2	70	-1.21

```
1 augment(fit2, newdata = temps, type.predict = "response")
```

	# A tibble: 2 × 2	temp .fitted
1	<dbl>	<dbl>
2	60	0.753
2	70	0.230

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

30

# A “Simple” Model Plot


As we'll see on the next slide, we will use the `augment` function to get the fitted probabilities into the original data, then plot.

- Note that we're just connecting the predictions made for observed `temp` values with `geom_line`, so the appearance of the function isn't as smooth as the actual logistic regression model.

## Plot our Model fit2

```
1 fit2_aug <- augment(fit2, type.predict = "response")
2
3 ggplot(fit2_aug, aes(x = temp, y = burst)) +
4   geom_point(alpha = 0.4) +
5   geom_line(aes(x = temp, y = .fitted),
6             col = "purple", size = 1.5) +
7   labs(title = "Fitted Logistic fit2 for Pr(burst)")
```



# Plot our Model fit2



432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

33

## Comparing fits of `fit1` and `fit2`



432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

34

# Try exponentiating `fit2` coefficients?

How can we interpret the coefficients of the model?

$$\text{logit}(burst) = \log(\text{odds}(burst))$$

## Exponentiating the slope is helpful

```
1 exp(-0.232)
```

```
[1] 0.7929461
```

## Exponentiating the slope helps

```
1 exp(-0.232)
```

```
[1] 0.7929461
```

Suppose Launch A's temperature was one degree higher than Launch B's.

- The **odds** of Launch A having a burst are 0.793 times as large as they are for Launch B.
- Odds Ratio estimate comparing two launches whose `temp` differs by 1 degree is 0.793

# Exponentiated and tidied slope `fit2`

```
1 tidy(fit2, exponentiate = TRUE, conf.int = TRUE, conf.level = 0.90) |>
2   filter(term == "temp") |>
3   gt() |> fmt_number(decimals = 3) |>
4   tab_options(table.font.size = 24)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
temp	0.793	0.108	-2.145	0.032	0.632	0.919

- What would it mean if the Odds Ratio for `temp` was 1?
- How about an odds ratio that was greater than 1?

# Standing Break

# Regression on a Binary Outcome

## Linear Probability Model (a linear model)

```
lm(event ~ predictor1 + predictor2 + ..., data = tiblename)
```

- $\Pr(\text{event})$  is linear in the predictors

## Logistic Regression Model (generalized linear model)

```
glm(event ~ pred1 + pred2 + ..., data = tiblename,  
    family = binomial(link = "logit"))
```

- Logistic Regression forces a prediction in  $(0, 1)$
- $\log(\text{odds}(\text{event}))$  is linear in the predictors

# The logistic regression model

$$\begin{aligned} \text{logit}(\text{event}) &= \log\left(\frac{\Pr(\text{event})}{1 - \Pr(\text{event})}\right) \\ \text{odds}(\text{event}) &= \frac{\Pr(\text{event})}{1 - \Pr(\text{event})} \\ \Pr(\text{event}) &= \frac{\text{odds}(\text{event}) + 1}{\exp(\text{logit}(\text{event}))} \\ \Pr(\text{event}) &= \frac{\exp(\text{logit}(\text{event}))}{1 + \exp(\text{logit}(\text{event}))} \end{aligned}$$

# model\_parameters() for fit2

```
1 model_parameters(fit2, ci = 0.90)
```

Parameter	Log-Odds	SE	90% CI	z	p
<hr/>					
(Intercept)	15.04	7.38	[ 4.95, 30.48]	2.04	0.041
temp	-0.23	0.11	[-0.46, -0.08]	-2.14	0.032

Uncertainty intervals (profile-likelihood) and p-values (two-tailed) computed using a Wald z-distribution approximation.

The model has a log- or logit-link. Consider using `exponentiate = TRUE` to interpret coefficients as ratios.

## Odds Ratios from model\_parameters()

```
1 model_parameters(fit2, exponentiate = TRUE, ci = 0.90)
```

Parameter	Odds Ratio	SE	90% CI	z	p
<hr/>					
(Intercept)	3.41e+06	2.52e+07	[141.15, 1.73e+13]	2.04	0.041
temp	0.79	0.09	[ 0.63, 0.92]	-2.14	0.032

## model fit2 slope (and CI)

Sample odds ratio for temp is 0.79, with 90% CI (0.63, 0.92)

- If launch 1 has a temperature 1 degree colder than launch 2, then our model estimates the odds of a burst to be 0.79 times as large (79% as large) for launch 2 as for launch 1.
- If our sample of launches was a random sample, then our 90% confidence interval suggests that if we generalize to the population of launches, then our data are consistent (at the 90% confidence level) with odds ratios between 0.63 and 0.92, assuming logistic regression assumptions are met.

# Compare `fit2` to a null model

- Likelihood Ratio test compares `fit2` to a model with only an intercept term (no `temp` information)

```
1 anova(fit2, test = "LRT")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: burst

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			22	28.267	
temp	1	7.952	21	20.315	0.004804 **
	---				
Signif. codes:	0	'***'	0.001	'**'	0.01 '*' 0.05 '.' 0.1 ' ' 1

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

43

## Other ANOVA options

- We can also get Rao's efficient score test (`test = "Rao"`) or Pearson's chi-square test (`test = "Chisq"`)

```
1 anova(fit2, test = "Rao")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: burst

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Rao	Pr(>Chi)
NULL			22	28.267		
temp	1	7.952	21	20.315	7.2312	0.007165 **
	---					
Signif. codes:	0	'***'	0.001	'**'	0.01 '*' 0.05 '.' 0.1 ' ' 1	

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

44

# Evaluating how well a logistic regression model predicts the outcome

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

45

## AUC and evaluating prediction quality

The Receiver Operating Characteristic (ROC) curve is the first approach we'll discuss today.

- Specifically, we will calculate the Area under this curve (sometimes labeled AUC or just C).

```
1 performance_roc(fit2)
```

AUC: 78.57%

- AUC falls between 0 and 1, and we interpret its result using the table on the next slide...

# Interpreting the AUC (C statistic)

## AUC Interpretation

0.5	A coin-flip. Model is no better than flipping a coin.
0.6	Still a fairly weak model.
0.7	Low end of an “OK” model fit.
0.8	Pretty good predictive performance.
0.9	Outstanding predictive performance.
1.0	Perfect predictive performance.

Recall that our `fit2` has AUC = 0.7857.

432 Class 03 | 2026-01-20 | <https://thomaselove.github.io/432-2026/>

47

## Classification Table (Confusion Matrix)

1. Select a decision rule.

- We’ll predict `burst = 1` if `fit2` model predicted probability > 0.5

2. Build a set of predictions, and make them a factor.

```
1 fit2_preds <- as_factor(ifelse(predict(fit2, type = "response") > 0.5, 1, 0))
```

3. Ensure the factor has “event occurs” first.

```
1 fit2_preds <- fct_relevel(fit2_preds, "1", "0")
```

# Classification Table

4. Obtain the actual “event” status, as a factor

```
1 fit2_actual <- fct_relevel(as_factor(orings1$burst), "1", "0")
```

5. Build the table

```
1 fit2_tab <- table(predicted = fit2_preds, actual = fit2_actual)
2 fit2_tab
```

	actual	
predicted	1	0
1	4	0
0	3	16

- Of the 4 launches predicted to have a burst, all 4 did.
- Of the 19 launches predicted to have no burst, 3 actually had a burst.

# Confusion Matrix Summaries

```
1 fit2_tab
```

	actual	
predicted	1	0
1	4	0
0	3	16

# A more complete Set of Summaries

```
1 confusionMatrix(fit2_tab)
```

Confusion Matrix and Statistics

	actual	
predicted	1	0
1	4	0
0	3	16

Accuracy : 0.8696  
95% CI : (0.6641, 0.9722)

No Information Rate : 0.6957  
P-Value [Acc > NIR] : 0.04928

Kappa : 0.6497

## Quality of Fit with `glance()` (1/2)

```
1 glance(fit2)
```

```
# A tibble: 1 × 8
  null.deviance df.null logLik   AIC   BIC deviance df.residual nobs
            <dbl>    <int>  <dbl> <dbl>  <dbl>      <dbl>     <int> <int>
1        28.3       22  -10.2  24.3  26.6     20.3        21     23
```

- `nobs` = we fit `fit2` using 23 observations
- null model (intercept) has 22 residual df (`df.null`) with `null.deviance` of 28.3
- `fit2` (includes `temp`) has 21 residual df (`df.residual`) with `deviance` of 20.3
  - The deviance quantifies what the model **doesn't** explain

# Quality of Fit with `glance()` (2/2)

```
1 glance(fit2) |>
2   gt() |>
3   fmt_number(columns = c(-df.null, -df.residual, -nobs), decimals = 2) |>
4   tab_options(table.font.size = 24)
```

null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
28.27	22	-10.16	24.32	26.59	20.32	21	23

- Our `fit2` has `deviance` =  $-2 \times \log \text{likelihood}$  (`logLik`)
- `AIC` and `BIC` are for comparing models for the same outcome, as in linear regression (smaller values indicate better fits, as usual.)

# `model_performance(fit2)` (1/5)

```
1 model_performance(fit2)

# Indices of model performance

AIC | AICc | BIC | Tjur's R2 | RMSE | Sigma | Log_loss | Score_log
-----
```

24.3		24.9		26.6		0.338		0.372		1		0.442		-2.957
------	--	------	--	------	--	-------	--	-------	--	---	--	-------	--	--------

```
AIC | Score_spherical | PCP
-----
```

24.3		0.149		0.720
------	--	-------	--	-------

- `AIC` and `BIC` are Akaike and Bayes information criteria
- `AICc` is a corrected AIC (correction for small sample size)
- `Sigma` is the estimated residual standard deviation
- `RMSE` estimates the root mean squared error

## model\_performance(fit2) (2/5)

Tjur's R2 = 0.338 for fit2

- Tjur's  $R^2$  is Tjur's coefficient of determination. Higher values indicate better fit.
  - Other choices: Cox-Snell , Nagelkerke , McFadden .
  - These pseudo- measures do **some** of what  $R^2$  does in linear regression are available in logistic regression.
  - Pseudo- measures don't describe proportionate reduction in error.
- Tjur's  $R^2$  can be calculated as follows:
  - For each level of the dependent variable, find the mean of the predicted probabilities of an event.
  - Take the absolute value of the difference between these means.

## model\_performance(fit2) (3/5)

Log\_loss = 0.442 for fit2

- Log\_loss quantifies prediction quality. If  $y_i$  is the actual/true value (1 or 0),  $p_i$  is the predicted probability, and  $\ln$  is the natural logarithm, then

$$\text{Log\_loss}_i = -[y_i \ln(p_i) + (1 -$$

- Model Log\_loss = sum of individual Log\_loss values.
- Lower Log\_loss values indicate better predictions.

## model\_performance() (4/5)

Score_log		Score_spherical
-2.957		0.149

- `Score_log` and `Score_spherical` are two other scoring rules for predictive performance in a logistic regression.
- `Score_log` takes values from  $[-, 0]$  with values closer to 0 indicating a more accurate model.
- `Score_spherical` takes values from  $[0, 1]$  with values closer to 1 indicating a more accurate model.
- See [this link](#) for more details.

## model\_performance() (5/5)


$$\text{PCP} = 0.720 \text{ for fit2}$$

- `PCP` is called the percentage of correct predictions
- `PCP` = sum of predicted probabilities where  $y=1$ , plus the sum of  $1 - \text{predicted probabilities}$  where  $y=0$ , divided by the number of observations
  - `PCP` ranges from 0 (worst) to 1 (best).
  - In general, the PCP should exceed 0.5.
- See [this link](#) for more details.


# Checking the fit2 model

```
1 check_model(fit2, check = c("pp_check", "outliers"))
```

Posterior Predictive Check  
Model-predicted intervals should include observed data points




Influential Observations  
Points should be inside the contour lines



# Checking the fit2 model

```
1 check_model(fit2, check = c("binned_residuals", "qq"))
```

Binned Residuals  
Points should be within error bounds



# Coming Up...

- The next several classes will be dedicated to providing more examples and more tools for working with linear regression and with logistic regression models.
- You should now have everything you need to do Lab 2.