

432 Class 05

Thomas E. Love, Ph.D.

2026-01-27

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

1

Today's Agenda

- The HELP trial, again
- Incorporating Non-Linearity into our models
 - Polynomial terms
 - Restricted Cubic Splines

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

2

Today's R Setup

```
1 knitr::opts_chunk$set(comment = NA)
2
3 library(janitor)
4 library(naniar)
5 library(broom); library(gt); library(patchwork)
6
7 library(haven)           ## for zapping labels
8 library(mosaic)          ## auto-loads mosaicData - data source
9
10 library(rms)             ## auto-loads Hmisc
11 library(easystats)
12 library(tidyverse)
13
14 theme_set(theme_bw())
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

3

Reminders: The HELP Study

Health Evaluation and Linkage to Primary Care (HELP) was a clinical trial of adult inpatients recruited from a detoxification unit.

- We have baseline data for each subject on several variables, including two outcomes:

Variable	Description
----------	-------------

cesd	Center for Epidemiologic Studies-Depression
------	---

cesd_hi	cesd above 15 (indicates high risk)
---------	-------------------------------------

Potential Predictors in `help1`

Variable	Description
<code>age</code>	subject age (in years)
<code>sex</code>	female (n = 107) or male (n = 346)
<code>subst</code>	substance abused (alcohol, cocaine, heroin)
<code>mcs</code>	SF-36 Mental Component Score
<code>pcs</code>	SF-36 Physical Component Score
<code>pss_fr</code>	perceived social support by friends

- See <https://nhorton.people.amherst.edu/help/> for more.

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

5

`help1` data load

```
1 help1 <- tibble(mosaicData::HELPrc) |>
2   select(id, cesd, age, sex, subst = substance, mcs, pcs, pss_fr) |>
3   zap_label() |>
4   mutate(across(where(is.character), as_factor),
5         id = as.character(id),
6         cesd_hi = factor(as.numeric(cesd >= 16)))
7
8 dim(help1); n_miss(help1)

[1] 453   9
[1] 0

1 head(help1, 5)

# A tibble: 5 × 9
  id    cesd    age sex    subst      mcs    pcs pss_fr cesd_hi
  <chr> <int> <int> <fct>    <fct>    <dbl> <dbl> <int> <fct>
1 1       49     37 male  cocaine  25.1    58.4     0  1
2 2       30     37 male  alcohol  26.7    36.0     1  1
3 3       39     26 male  heroin   6.76    74.8    13  1
4 4       15     39 female heroin  44.0    61.9    11  0
5 5       39     32 male  cocaine  21.7    37.3    10  1
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

6

Can we use `pcs` to predict `cesd`?

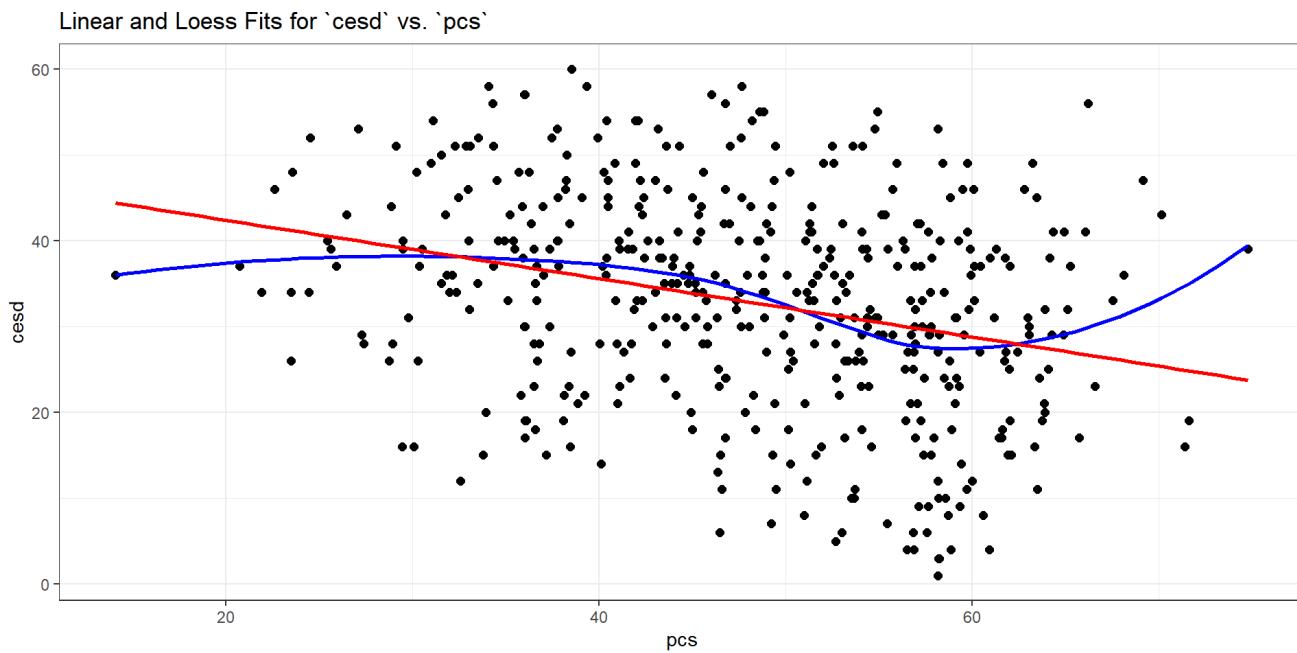
Does the `loess` smooth match up well with the linear fit?

```
1 ggplot(help1, aes(x = pcs, y = cesd)) +  
2   geom_point(size = 2) +  
3   geom_smooth(method = "loess", formula = y ~ x, se = FALSE, col = "blue") +  
4   geom_smooth(method = "lm", formula = y ~ x, se = FALSE, col = "red") +  
5   labs(title = "Linear and Loess Fits for `cesd` vs. `pcs`")
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

7

Can we use `pcs` to predict `cesd`?



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

8

A simple linear regression: fitA

```
1 dd <- datadist(help1); options(datadist = "dd")
2
3 fitA <- ols(cesd ~ pcs, data = help1, x = TRUE, y = TRUE)
4
5 fitA$coefficients
```

	Intercept	pcs
	49.1673458	-0.3396495

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

9

Our simple linear regression

```
1 fitA
```

Linear Regression Model

```
ols(formula = cesd ~ pcs, data = help1, x = TRUE, y = TRUE)
```

	Model	Likelihood	Discrimination
		Ratio Test	Indexes
Obs	453	LR chi2	40.57
sigma1	1.9796	d.f.	1
d.f.	451	Pr(> chi2)	0.0000
		R2	0.086
		R2 adj	0.084
		g	4.177

Residuals

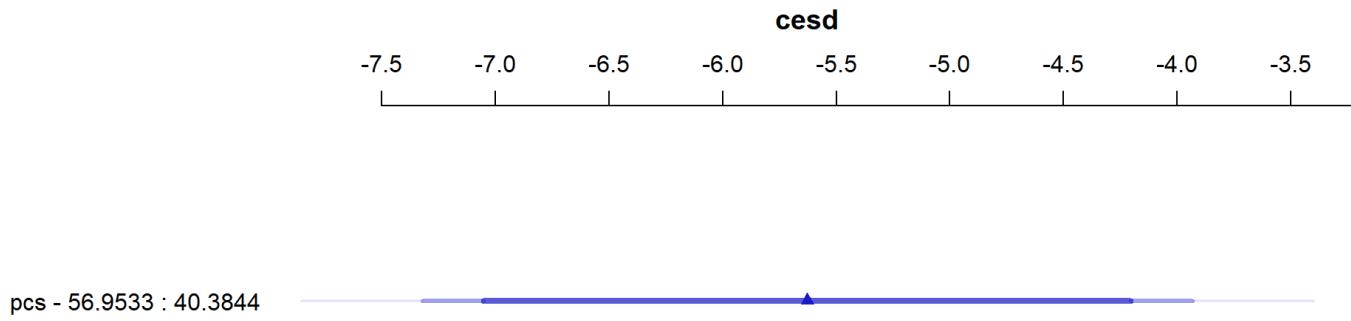
	Min	1Q	Median	3Q	Max
	-28.4116	-7.8036	0.6846	8.7917	29.3281

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

10

Effect Sizes in `fitA`

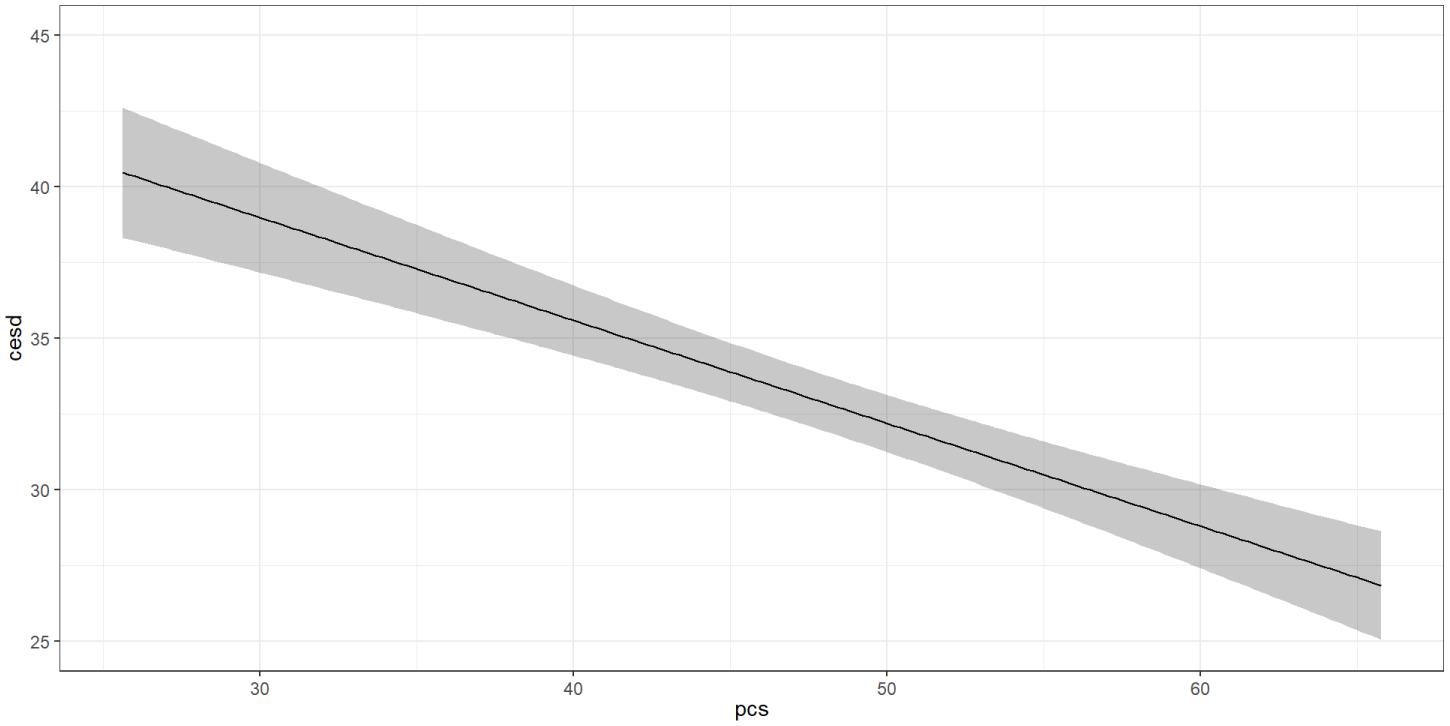
```
1 plot(summary(fitA))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

11

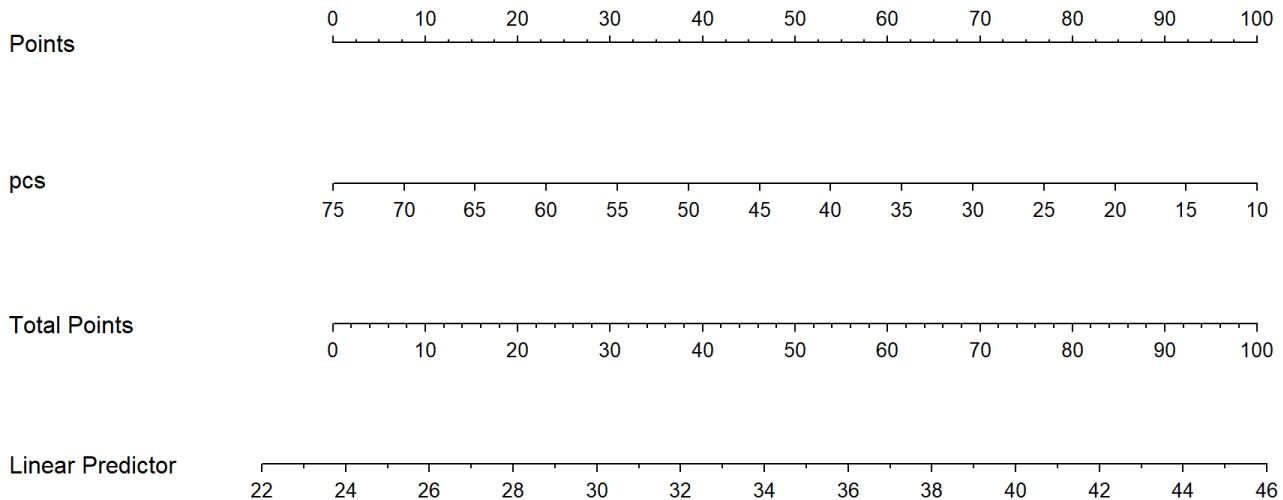
```
1 ggplot(Predict(fitA, conf.int = 0.90))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

12

```
1 plot(nomogram(fitA))
```



Using `ols` to fit a larger model

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fitB <- ols(cesd ~ pcs + subst + pss_fr + sex,
5               data = help1, x = TRUE, y = TRUE)
6
7 fitB$coefficients
```

	Intercept	pcs	subst=cocaine	subst=heroin	pss_fr
sex=male	53.7511151	-0.2574023	-3.8664109	0.2322071	-0.5370221
	-4.8446977				

- Can use `model_parameters()` and `model_performance()` with `fitB` or other `ols()` fits.
- We could also fit this model, naturally, using `lm()` instead.

Contents of fitB?

```
1 fitB
```

Linear Regression Model

```
ols(formula = cesd ~ pcs + subst + pss_fr + sex, data = help1,  
x = TRUE, y = TRUE)
```

	Model Likelihood	Discrimination			
	Ratio Test	Indexes			
Obs	453	LR chi2	76.43	R2	0.155
sigma11	5.662	d.f.	5	R2 adj	0.146
d.f.	447	Pr(> chi2)	0.0000	g	5.625

Residuals

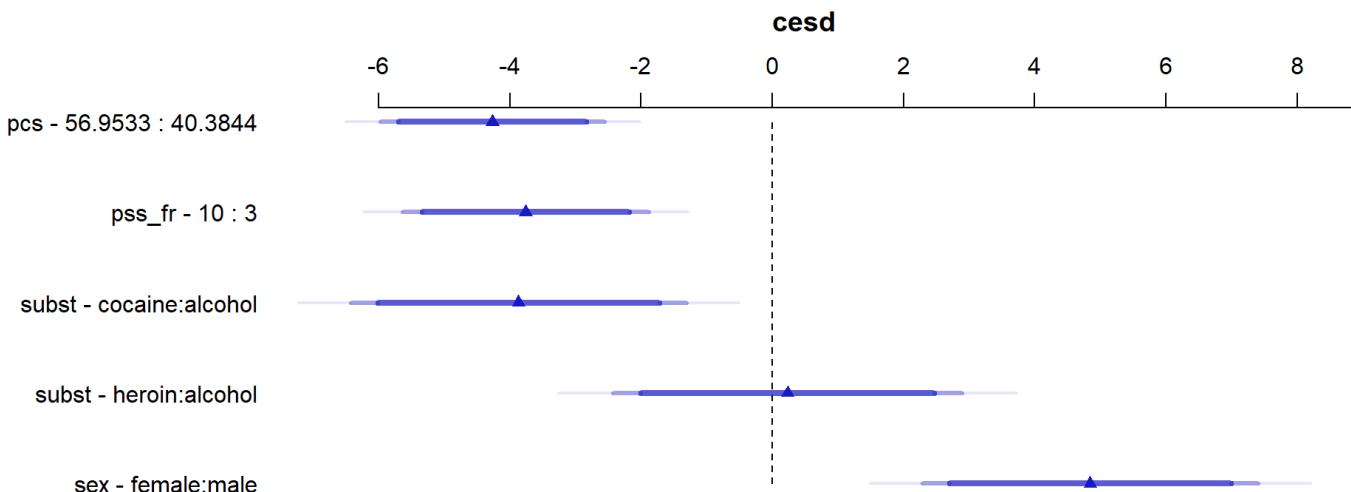
Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

15

Effect Sizes in fitB

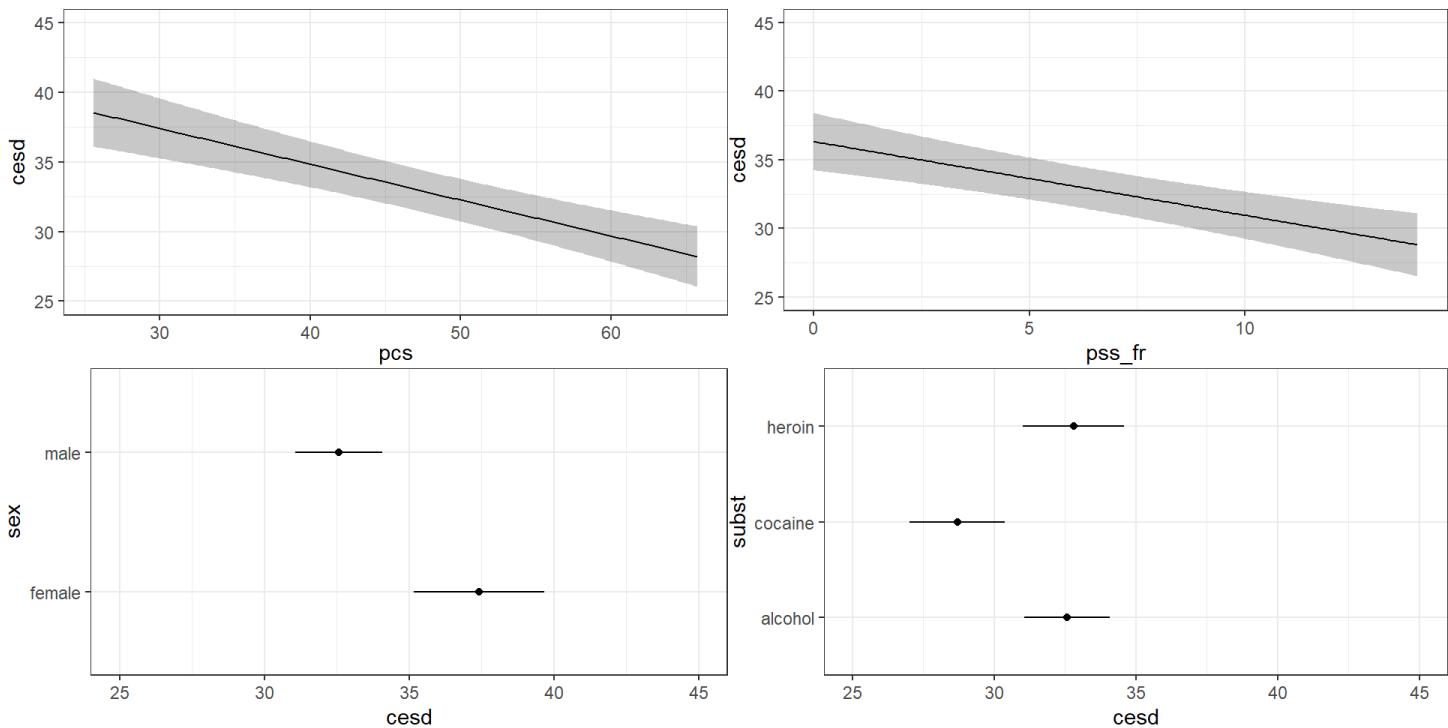
```
1 plot(summary(fitB))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

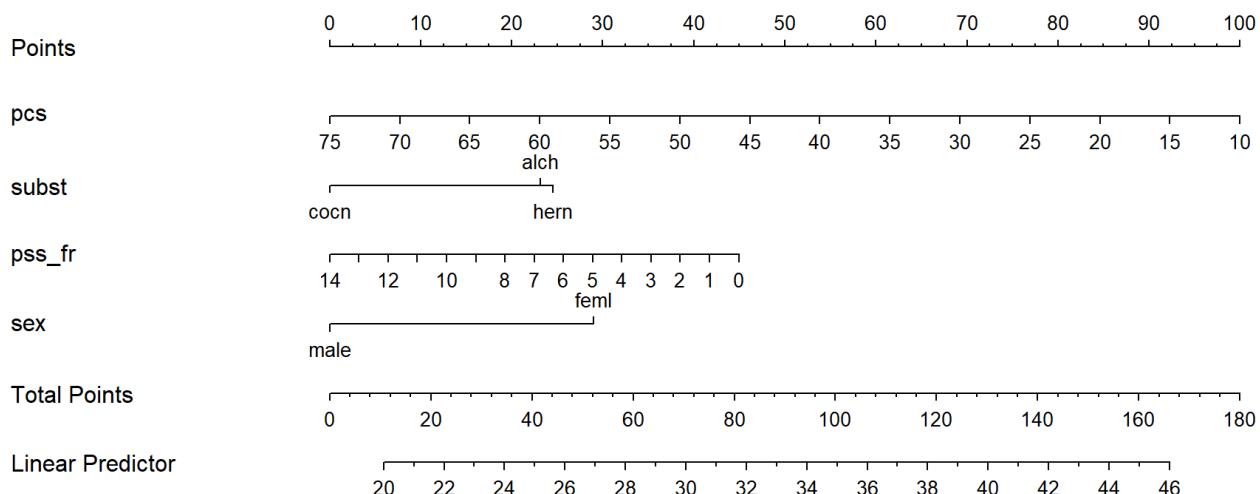
16

```
1 ggplot(Predict(fitB, conf.int = 0.90))
```



A Nomogram for fitB

```
1 plot(nomogram(fitB, abbrev = TRUE))
```



Non-Linear Terms

In building a linear regression model, we're most often going to be thinking about:

- for quantitative predictors, some curvature...
 - perhaps polynomial terms
 - but more often restricted cubic splines
- for any predictors, possible interactions
 - between categorical predictors
 - between categorical and quantitative predictors
 - between quantitative predictors

Non-Linear Terms: Polynomials

Polynomial Regression

A polynomial in the variable x of degree D is a linear combination of the powers of x up to D. For example:

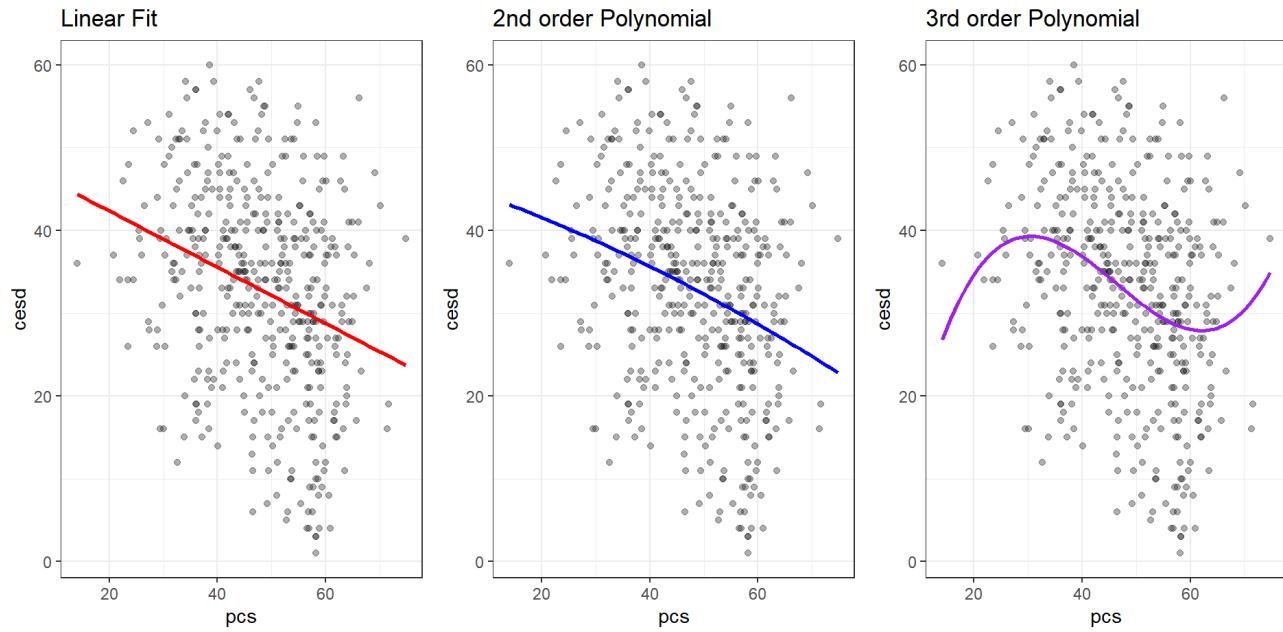
- Linear: $y = \beta_0 + \beta_1 x$
- Quadratic: $y = \beta_0 + \beta_1 x + \beta_2 x^2$
- Cubic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$
- Quartic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4$

Fitting such a model creates a **polynomial regression**.

Plotting the Polynomials

```
1 p1 <- ggplot(help1, aes(x = pcs, y = cesd)) +
2   geom_point(alpha = 0.3) +
3   geom_smooth(formula = y ~ x, method = "lm",
4             col = "red", se = FALSE) +
5   labs(title = "Linear Fit")
6
7 p2 <- ggplot(help1, aes(x = pcs, y = cesd)) +
8   geom_point(alpha = 0.3) +
9   geom_smooth(formula = y ~ poly(x, 2), method = "lm",
10            col = "blue", se = FALSE) +
11   labs(title = "2nd order Polynomial")
12
13 p3 <- ggplot(help1, aes(x = pcs, y = cesd)) +
14   geom_point(alpha = 0.3) +
15   geom_smooth(formula = y ~ poly(x, 3), method = "lm",
16             col = "purple", se = FALSE) +
17   labs(title = "3rd order Polynomial")
18
```

Plotting the Polynomials



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

23

Adding a polynomial in `pcs`

Can we predict `cesd` with a polynomial in `pcs`?

Yes, with `ols()` and `pol()`, as follows:

```
fitA <- ols(cesd ~ pcs, data = help1, x = TRUE, y = TRUE)
fitA_2 <- ols(cesd ~ pol(pcs,2), data = help1, x = TRUE, y = TRUE)
fitA_3 <- ols(cesd ~ pol(pcs,3), data = help1, x = TRUE, y = TRUE)
```

With `lm()`, we use `poly()` instead of `pol()`...

```
lmfitA <- lm(cesd ~ pcs, data = help1)
lmfitA_2 <- lm(cesd ~ poly(pcs,2), data = help1)
lmfitA_3 <- lm(cesd ~ poly(pcs,3), data = help1)
```

Raw vs. Orthogonal Polynomials

Predict `cesd` using `pcs` with a “raw polynomial of degree 2.”

```
1 (temp1 <- lm(cesd ~ pcs + I(pcs^2), data = help1))
```

Call:

```
lm(formula = cesd ~ pcs + I(pcs^2), data = help1)
```

Coefficients:

(Intercept)	pcs	I(pcs^2)
46.400713	-0.213627	-0.001356

Predicted `cesd` for `pcs` = 40 is

```
cesd = 46.400713 - 0.213627 (40) - 0.001356 (40^2)
      = 46.400713 - 8.545080 - 2.169600
      = 35.686
```

Does the raw polynomial match our expectations?

```
1 temp1 <- lm(cesd ~ pcs + I(pcs^2), data = help1)
2
3 augment(temp1, newdata = tibble(pcs = 40)) |>
4   gt() |> tab_options(table.font.size = 24)
```

pcs	.fitted
40	35.6856

This matches our “by hand” calculation.

- But it turns out most regression models use *orthogonal* rather than raw polynomials...

Fitting an Orthogonal Polynomial

Predict `cesd` using `pcs` with an *orthogonal* polynomial of degree 2.

```
1 (temp2 <- lm(cesd ~ poly(pcs, 2), data = help1))
```

Call:

```
lm(formula = cesd ~ poly(pcs, 2), data = help1)
```

Coefficients:

(Intercept)	poly(pcs, 2)1	poly(pcs, 2)2
32.848	-77.876	-3.944

This looks very different from our previous version of the model. What happens when we make a prediction, though?

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

27

Orthogonal Polynomial Model Prediction

Remember that in our raw polynomial model, our “by hand” and “using R” calculations each predicted `cesd` for a subject with `pcs` = 40 to be 35.686.

What happens with the orthogonal polynomial model `temp2`?

```
1 augment(temp2, newdata = data.frame(pcs = 40)) |>
2   gt() |> tab_options(table.font.size = 24)
```

pcs	.fitted
40	35.6856

- No change in the prediction.

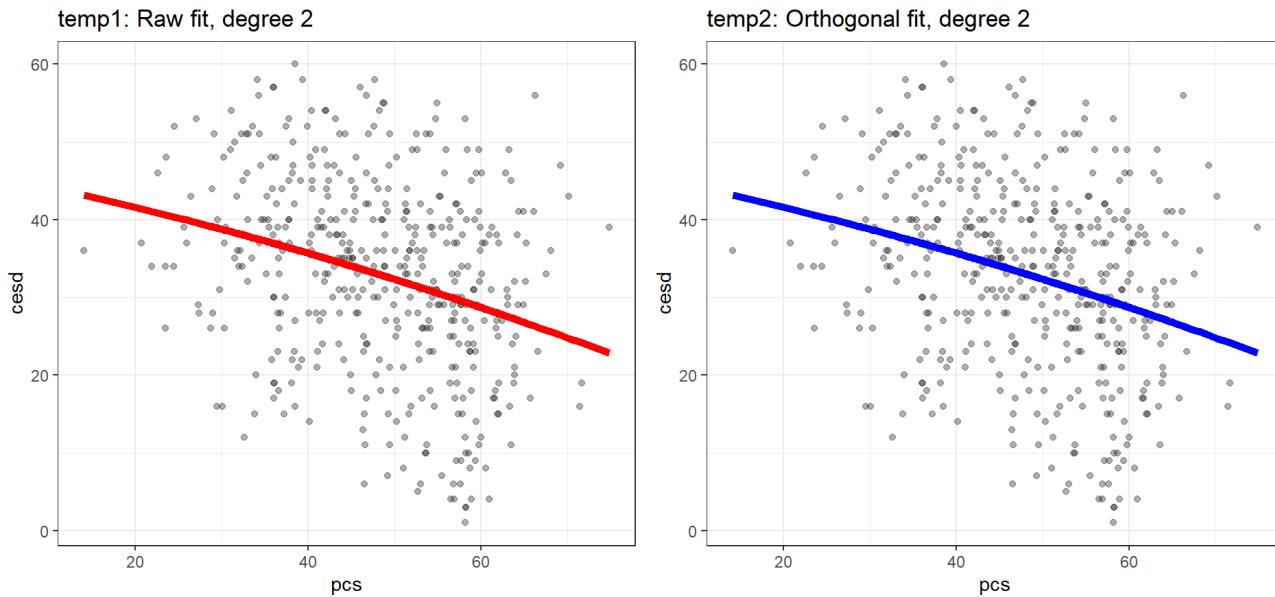
Fits of raw vs orthogonal polynomials

```
1 temp1_aug <- augment(temp1, help1)
2 temp2_aug <- augment(temp2, help1)
3
4 p1 <- ggplot(temp1_aug, aes(x = pcs, y = cesd)) +
5   geom_point(alpha = 0.3) +
6   geom_line(aes(x = pcs, y = .fitted), col = "red", linewidth = 2) +
7   labs(title = "temp1: Raw fit, degree 2")
8
9 p2 <- ggplot(temp2_aug, aes(x = pcs, y = cesd)) +
10  geom_point(alpha = 0.3) +
11  geom_line(aes(x = pcs, y = .fitted), col = "blue", linewidth = 2) +
12  labs(title = "temp2: Orthogonal fit, degree 2")
13
14 p1 + p2 +
15   plot_annotation(title = "Comparing Two Methods of Fitting a Quadratic Polynomia
```

- The two models are, in fact, identical.

Fits of raw vs orthogonal polynomials

Comparing Two Methods of Fitting a Quadratic Polynomial



Why use orthogonal polynomials?

- The main reason is to avoid having to include powers of our predictor that are highly collinear.
- Variance Inflation Factor assesses collinearity...

```
1 rms::vif(temp1)      ## from rms package  
pcs I(pcs^2)  
54.66793 54.66793
```

- Orthogonal polynomial terms are uncorrelated...

```
1 rms::vif(temp2)  
poly(pcs, 2)1 poly(pcs, 2)2  
1           1
```

Why orthogonal polynomials?

An **orthogonal polynomial** sets up a model design matrix and then scales those columns so that each column is uncorrelated with the others. The tradeoff is that the raw polynomial is a lot easier to explain in terms of a single equation in the simplest case.

Actually, we'll often use splines instead of polynomials, which are more flexible and require less maintenance, but at the cost of pretty much requiring you to focus on visualizing their predictions rather than their equations.

fitA with a cubic polynomial

```
1 dd <- datadist(help1); options(datadist = "dd")
2
3 fitA_3 <- ols(cesd ~ pol(pcs,3), data = help1, x = TRUE, y = TRUE)
4
5 fitA_3$coefficients
```

	Intercept	pcs	pcs^2	pcs^3
-1.340758e+01	4.132348e+00	-1.009667e-01	7.268386e-04	

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

33

Our model fitA_3

```
1 fitA_3
```

Linear Regression Model

```
ols(formula = cesd ~ pol(pcs, 3), data = help1, x = TRUE, y = TRUE)
```

	Model Likelihood		Discrimination	
	Ratio	Test		Indexes
Obs	453	LR chi2	48.70	R2 0.102
sigma	1.8991	d.f.	3	R2 adj 0.096
d.f.	449	Pr(> chi2)	0.0000	g 4.556

Residuals

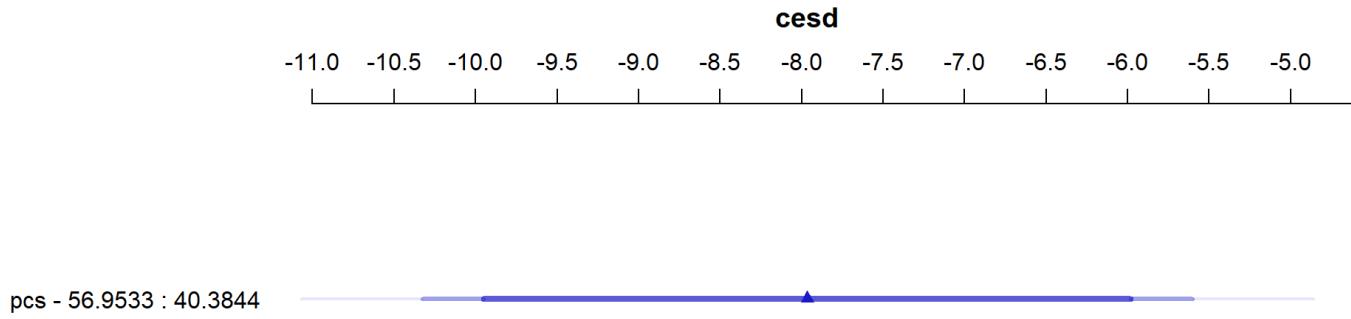
Min	1Q	Median	3Q	Max
-27.5245	-8.2651	0.7988	8.9004	27.4480

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

34

Effect Sizes in `fitA_3`

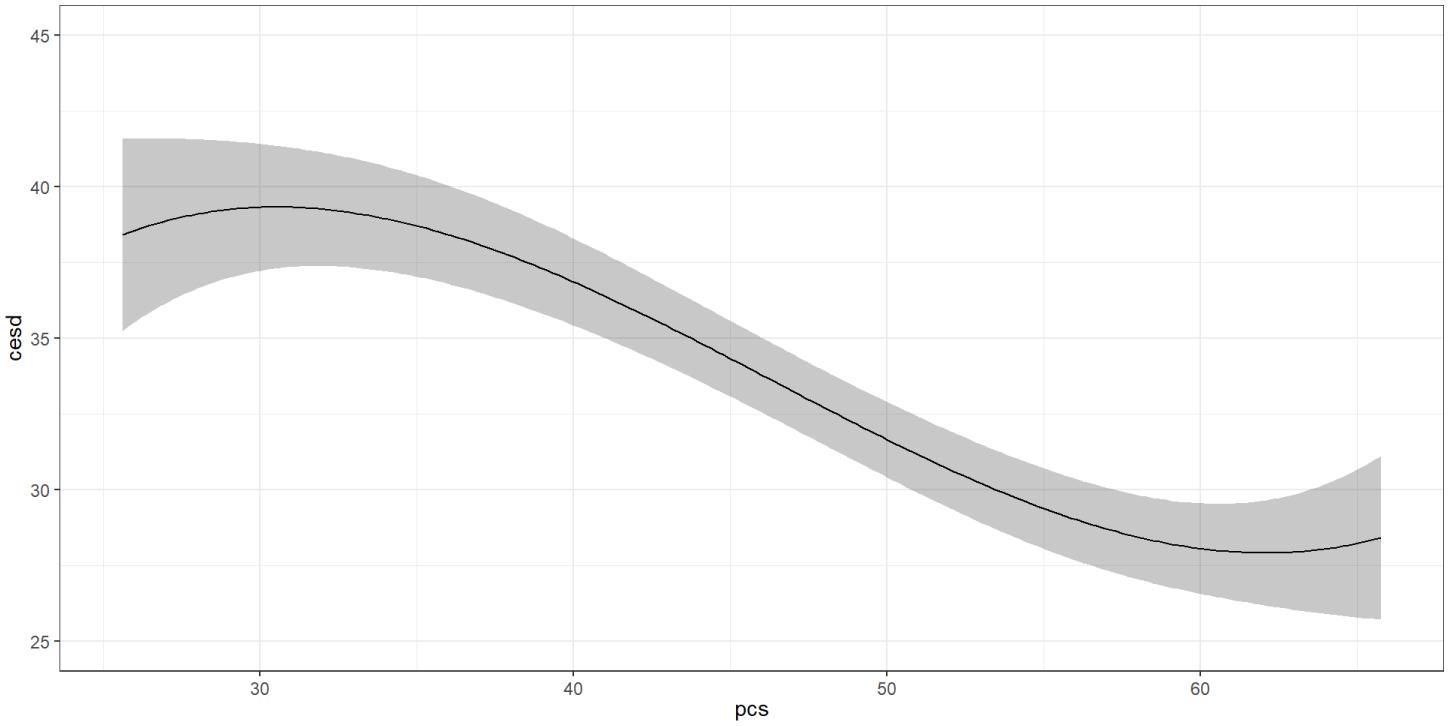
```
1 plot(summary(fitA_3))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

35

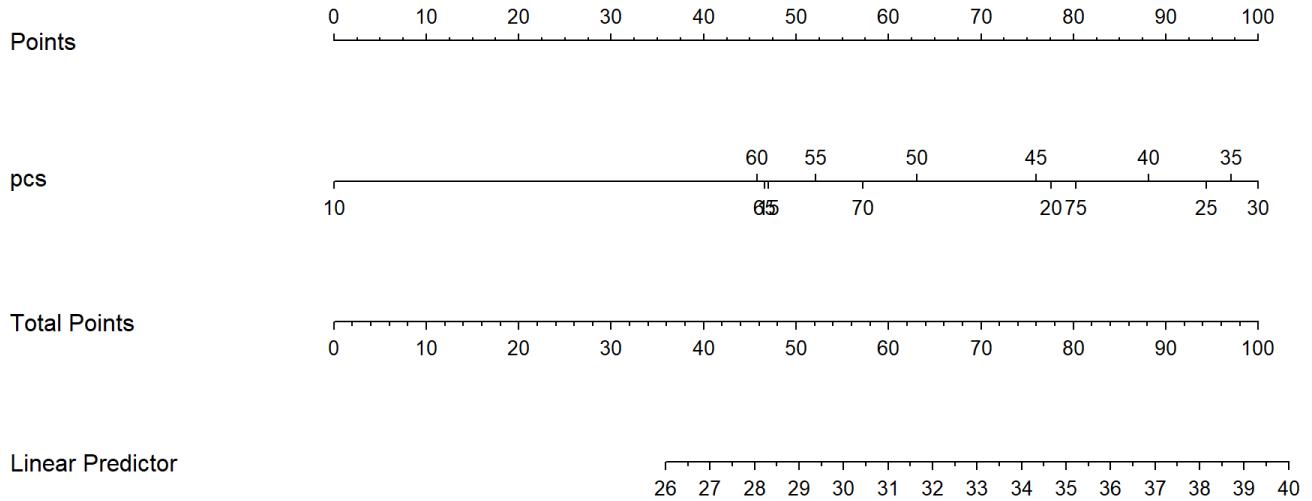
```
1 ggplot(Predict(fitA_3, conf.int = 0.90))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

36

```
1 plot(nomogram(fitA_3))
```



Fitting fitB including a polynomial

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fitB_3 <- ols(cesd ~ pol(pcs,3) + subst + pss_fr + sex,
5                 data = help1, x = TRUE, y = TRUE)
6
7 fitB_3$coefficients
```

	Intercept	pcs	pcs ²	pcs ³	subst=cocaine
subst=heroin	5.2983256376	3.2271532761	-0.0794837993	0.0005770243	-3.8581390102
		pss_fr	sex=male		
	0.0455051022	-0.5127744954	-4.5981834492		

Contents of fitB_3?

```
1 fitB_3
```

Linear Regression Model

```
ols(formula = cesd ~ pol(pcs, 3) + subst + pss_fr + sex, data = help1,  
x = TRUE, y = TRUE)
```

	Model Likelihood	Discrimination
	Ratio Test	Indexes
Obs	453	LR chi2 81.80
sigma11	5236	d.f. 7
d.f.	445	Pr(> chi2) 0.0000
		R2 0.165
		R2 adj 0.152
		g 5.808

Residuals

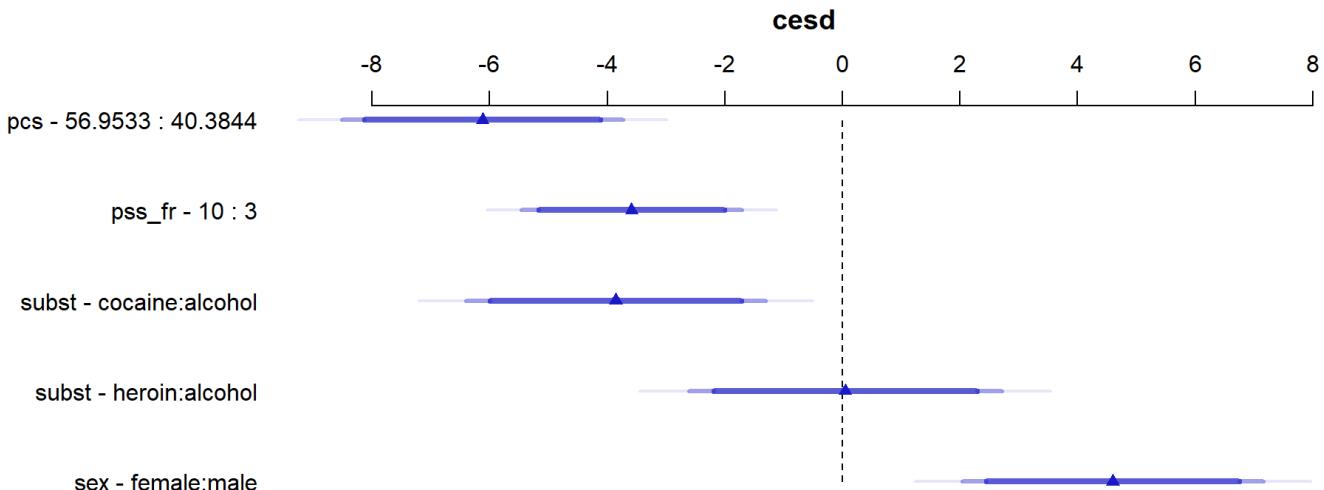
Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

39

Effect Sizes in fitB_3

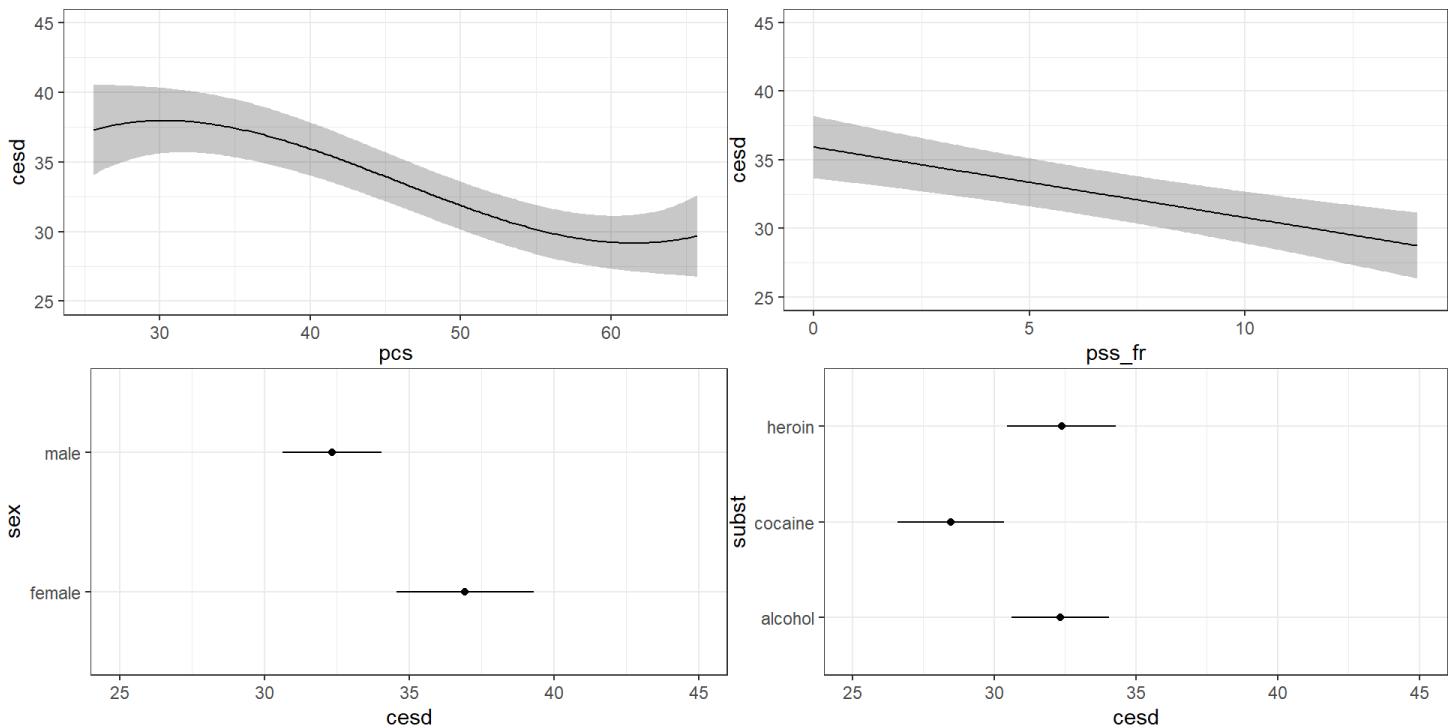
```
1 plot(summary(fitB_3))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

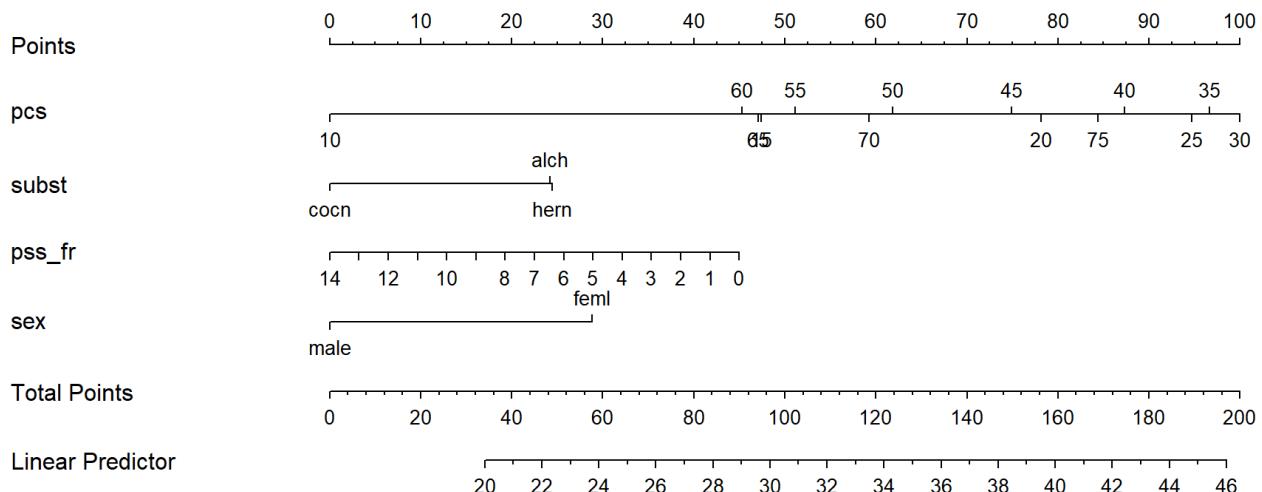
40

```
1 ggplot(Predict(fitB_3, conf.int = 0.90))
```



A Nomogram for fitB_3

```
1 plot(nomogram(fitB_3, abbrev = TRUE))
```



Standing Break

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

43

Non-Linear Terms: Splines

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

44

Types of Splines

- A **linear spline** is a continuous function formed by connecting points (called **knots** of the spline) by line segments.
- A **restricted cubic spline** is a way to build highly complicated curves into a regression equation in a fairly easily structured way.
- A restricted cubic spline is a series of polynomial functions joined together at the knots.
 - Such a spline gives us a way to flexibly account for non-linearity without over-fitting the model.

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

45

How complex should our spline be?

Restricted cubic splines can fit many types of non-linearity. Specifying the number of knots (usually 3, 4 or 5) is all you must do in R to get a reasonable result from a restricted cubic spline.

- 3 Knots, 2 degrees of freedom, allows the curve to “bend” once.
- 4 Knots, 3 degrees of freedom, lets the curve “bend” twice.
- 5 Knots, 4 degrees of freedom, lets the curve “bend” three times.

Restricted Cubic Splines with `ols`

Let's consider a restricted cubic spline model for `cesd` based on `pcs` with:

- 3 knots in `fitC3`, 4 knots in `fitC4`, and 5 knots in `fitC5`

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fitC3 <- ols(cesd ~ rcs(pcs, 3),
5                 data = help1, x = TRUE, y = TRUE)
6 fitC4 <- ols(cesd ~ rcs(pcs, 4),
7                 data = help1, x = TRUE, y = TRUE)
8 fitC5 <- ols(cesd ~ rcs(pcs, 5),
9                 data = help1, x = TRUE, y = TRUE)
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

47

Model `fitC3` (3-knot spline in `pcs`)

```
1 fitC3
```

Linear Regression Model

```
ols(formula = cesd ~ rcs(pcs, 3), data = help1, x = TRUE, y = TRUE)
```

	Model	Likelihood	Discrimination
		Ratio Test	Indexes
Obs	453	LR chi2	40.79
sigma11	9.901	d.f.	2
d.f.	450	Pr(> chi2)	0.0000
		R2	0.086
		R2 adj	0.082
		g	4.206

Residuals

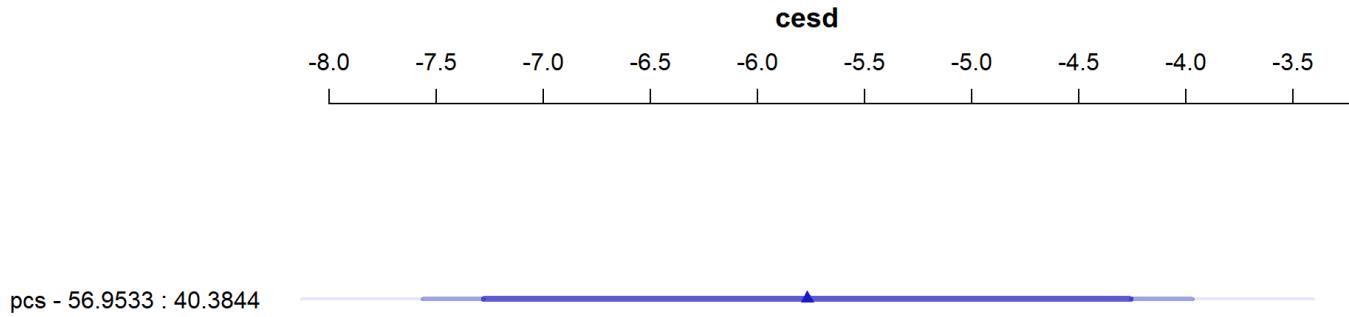
Min	1Q	Median	3Q	Max
-28.3462	-7.7005	0.5098	8.6376	29.8454

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

48

Effect Sizes in fitC3

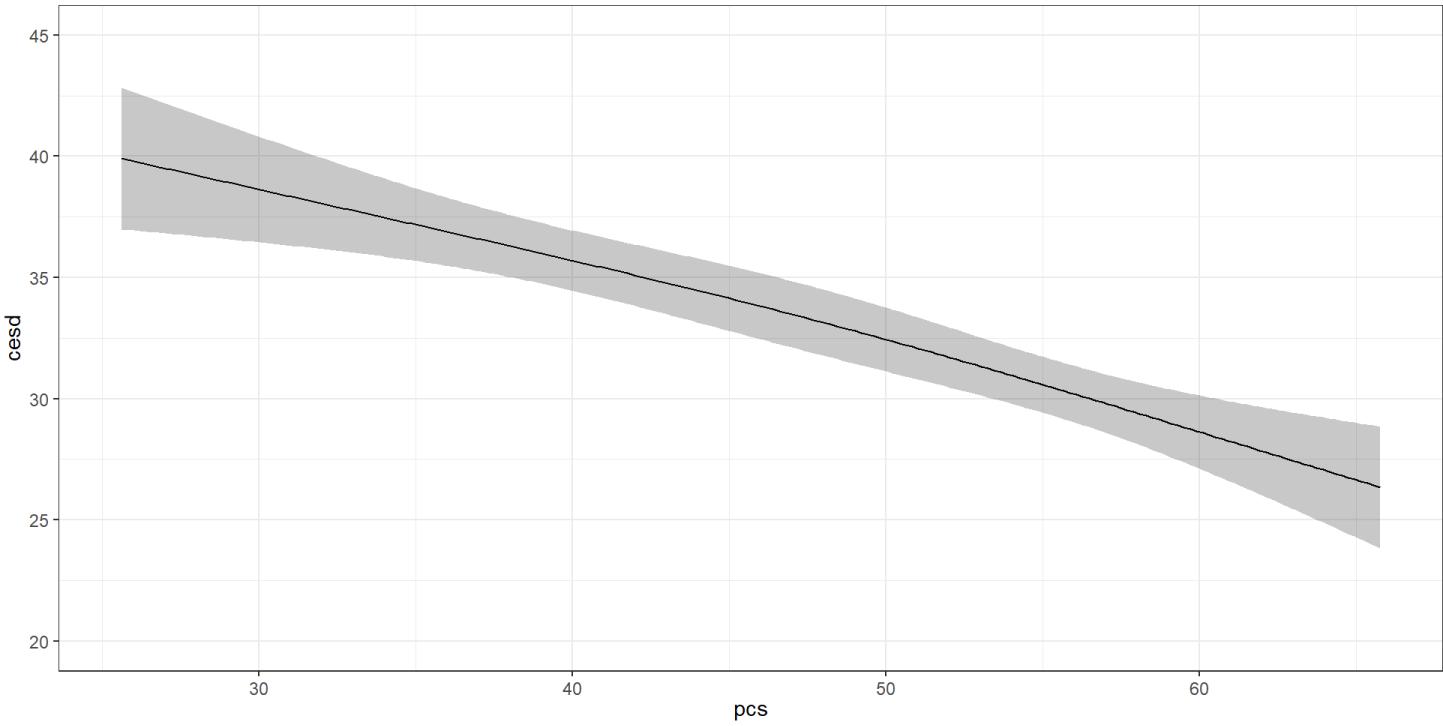
```
1 plot(summary(fitC3))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

49

```
1 ggplot(Predict(fitC3, conf.int = 0.90))
```

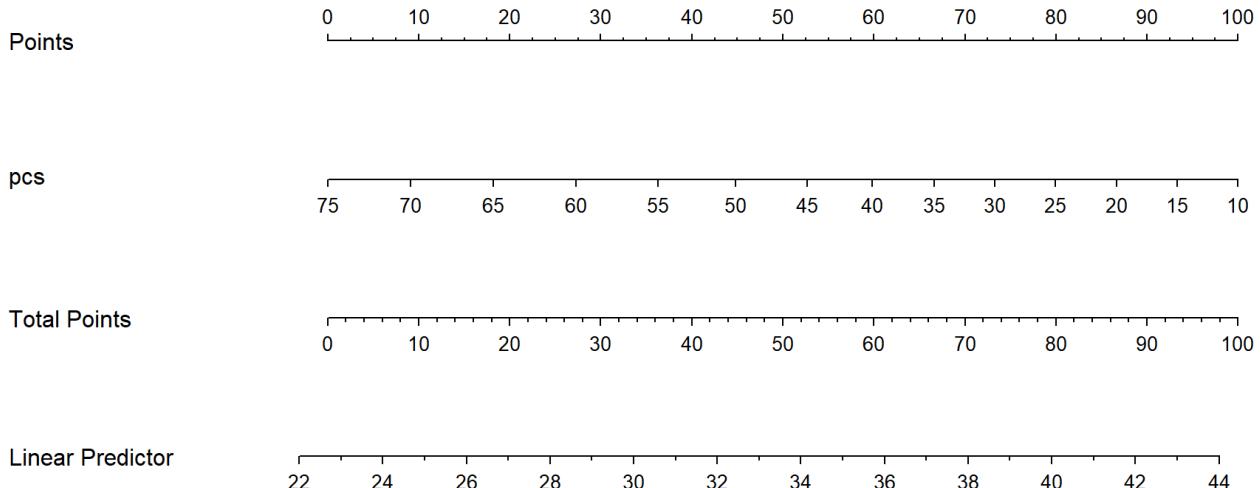


432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

50

A Nomogram for fitC3

```
1 plot(nomogram(fitC3, abbrev = TRUE))
```



Model fitC4 (4-knot spline in pcs)

```
1 fitC4
```

Linear Regression Model

```
ols(formula = cesd ~ rcs(pcs, 4), data = help1, x = TRUE, y = TRUE)
```

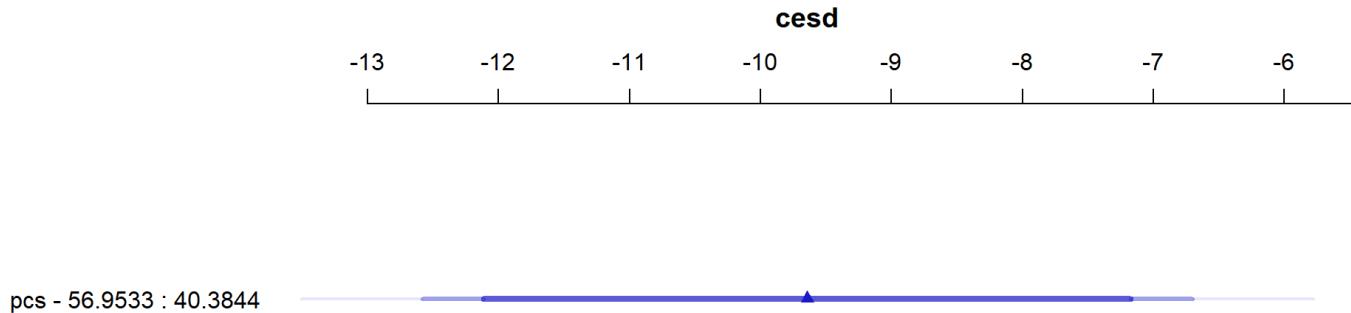
	Model	Likelihood	Discrimination
		Ratio Test	Indexes
Obs	453	LR chi2	51.31
sigma11	8648	d.f.	3
d.f.	449	Pr(> chi2)	0.0000
		R2	0.107
		R2 adj	0.101
		g	4.590

Residuals

Min	1Q	Median	3Q	Max
-28.3147	-8.2830	0.8559	8.8866	26.5458

Effect Sizes in fitC4

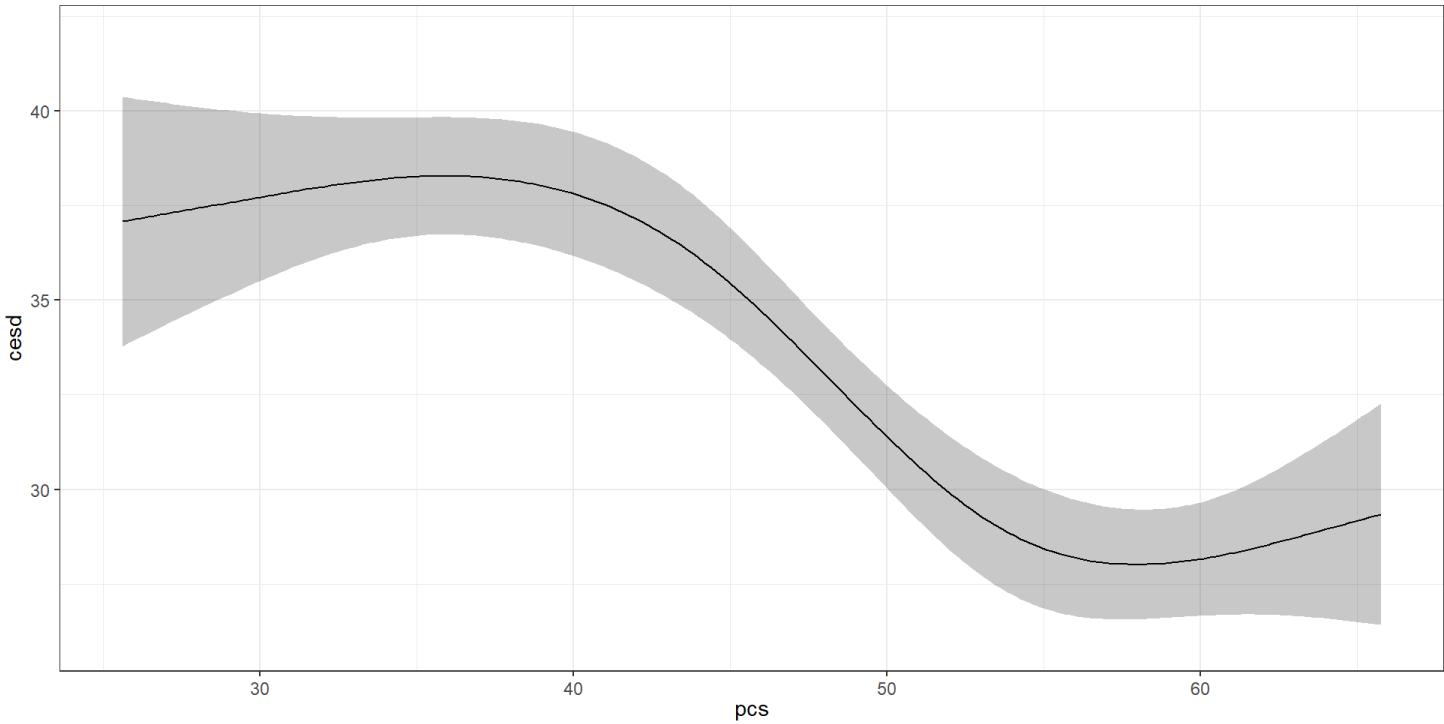
```
1 plot(summary(fitC4))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

53

```
1 ggplot(Predict(fitC4, conf.int = 0.90))
```

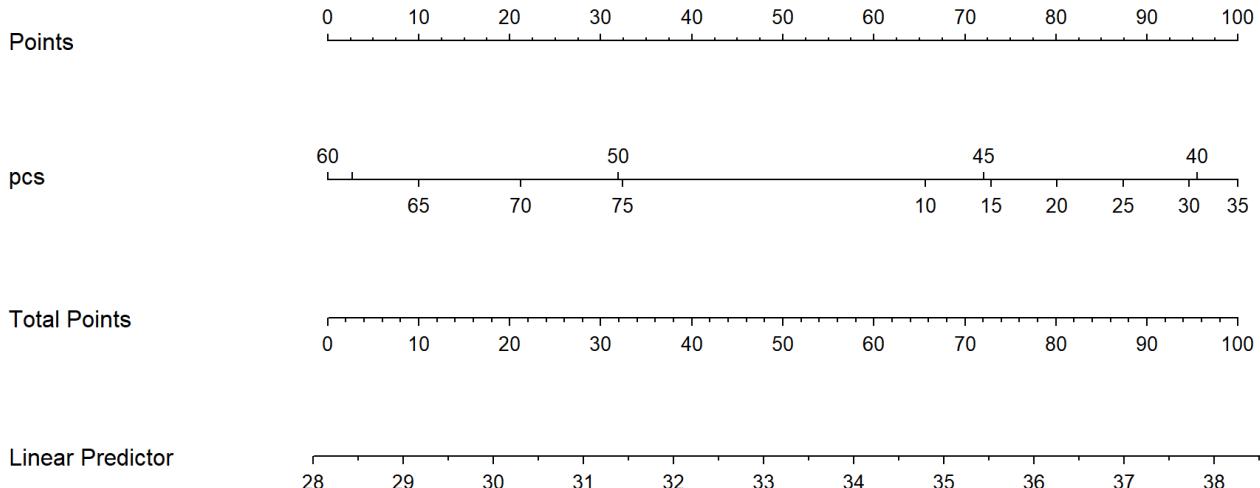


432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

54

A Nomogram for fitC4

```
1 plot(nomogram(fitC4, abbrev = TRUE))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

55

Model fitC5 (5-knot spline in pcs)

```
1 fitC5
```

Linear Regression Model

```
ols(formula = cesd ~ rcs(pcs, 5), data = help1, x = TRUE, y = TRUE)
```

	Obs	Model Likelihood		Discrimination	
		Ratio	Test	R2	Indexes
	453	LR	chi2	54.64	0.114
sigma11.8345		d.f.		4	R2 adj 0.106
d.f.	448	Pr(> chi2)	0.0000	g	4.744

Residuals

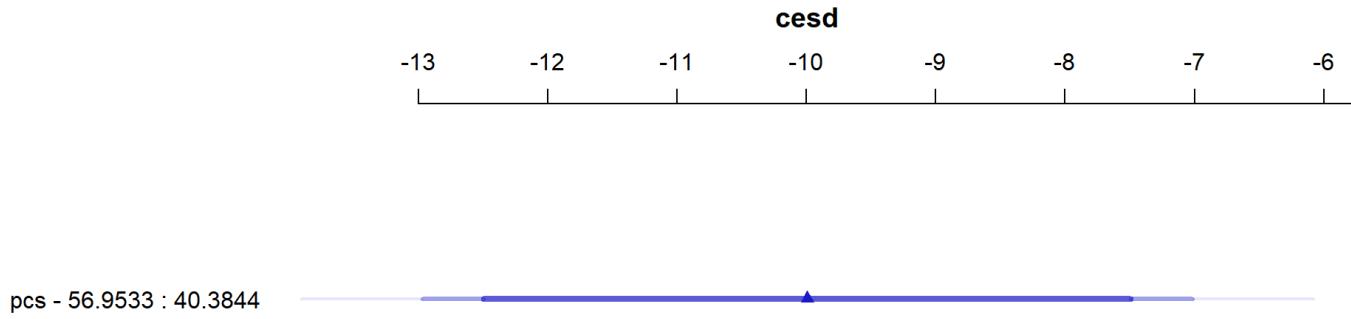
```
Min      1Q  Median      3Q      Max 
-29.396 -7.928  1.016   8.762  26.974
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

56

Effect Sizes in fitC5

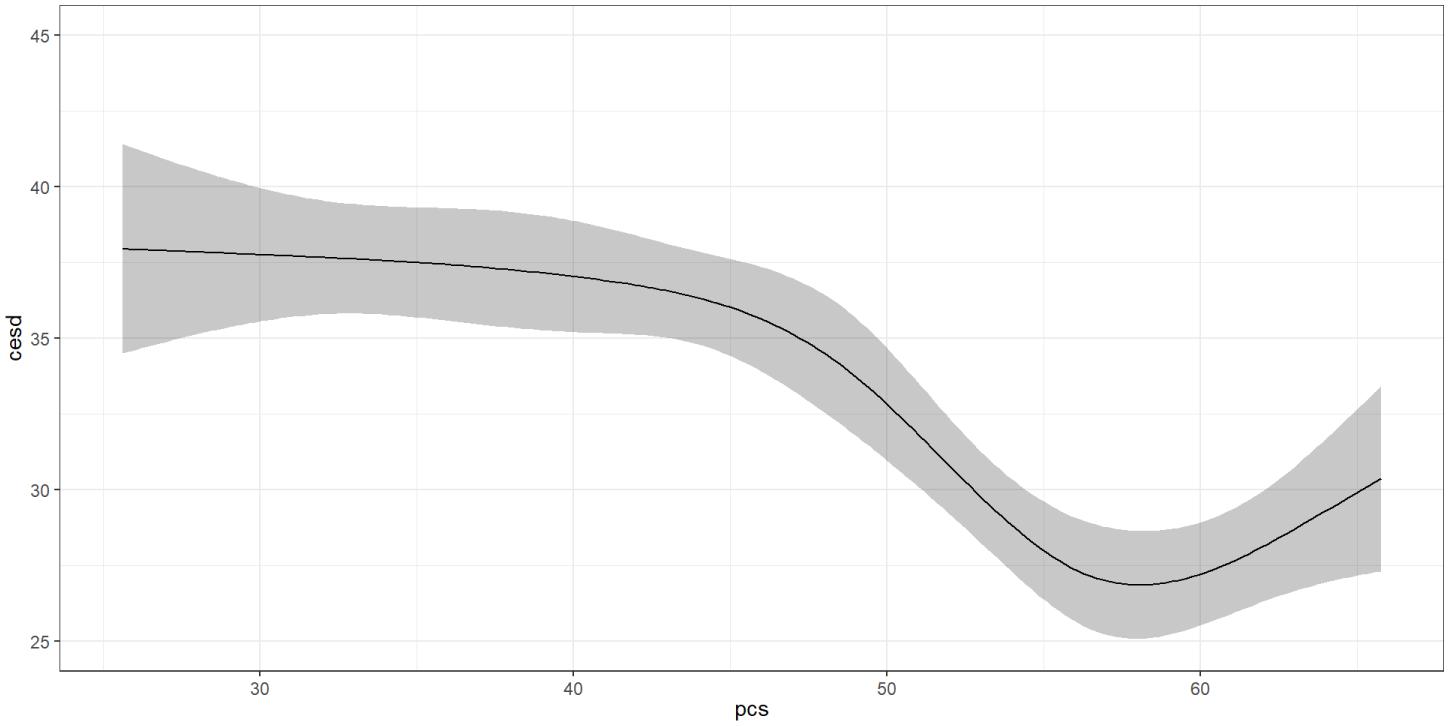
```
1 plot(summary(fitC5))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

57

```
1 ggplot(Predict(fitC5, conf.int = 0.90))
```

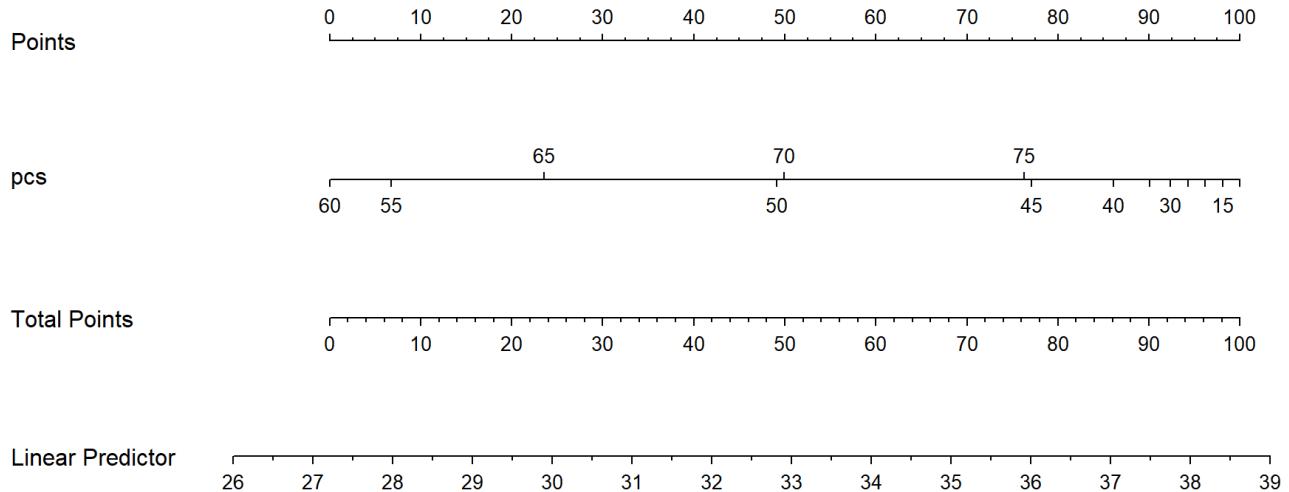


432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

58

A Nomogram for fitC5

```
1 plot(nomogram(fitC5, abbrev = TRUE))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

59

Fitting fitB including a 5-knot RCS

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fitB5 <- ols(cesd ~ rcs(pcs,5) + subst + pss_fr + sex,
5               data = help1, x = TRUE, y = TRUE)
6
7 fitB5$coefficients
```

	Intercept	pcs	pcs'	pcs''	pcs'''
	48.4263870	-0.1151961	0.2519557	-4.9683530	15.6487400
subst=cocaine	subst=heroin		pss_fr	sex=male	
	-3.7580871	-0.1528491	-0.4985050	-4.7527659	

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

60

Contents of fitB5?

```
1 fitB5
```

Linear Regression Model

```
ols(formula = cesd ~ rcs(pcs, 5) + subst + pss_fr + sex, data = help1,  
x = TRUE, y = TRUE)
```

	Model Likelihood	Discrimination
	Ratio Test	Indexes
Obs	453	LR chi2 86.99
sigma11	4.707	d.f. 8
d.f.	444	Pr(> chi2) 0.0000
		R2 0.175
		R2 adj 0.160
		g 5.991

Residuals

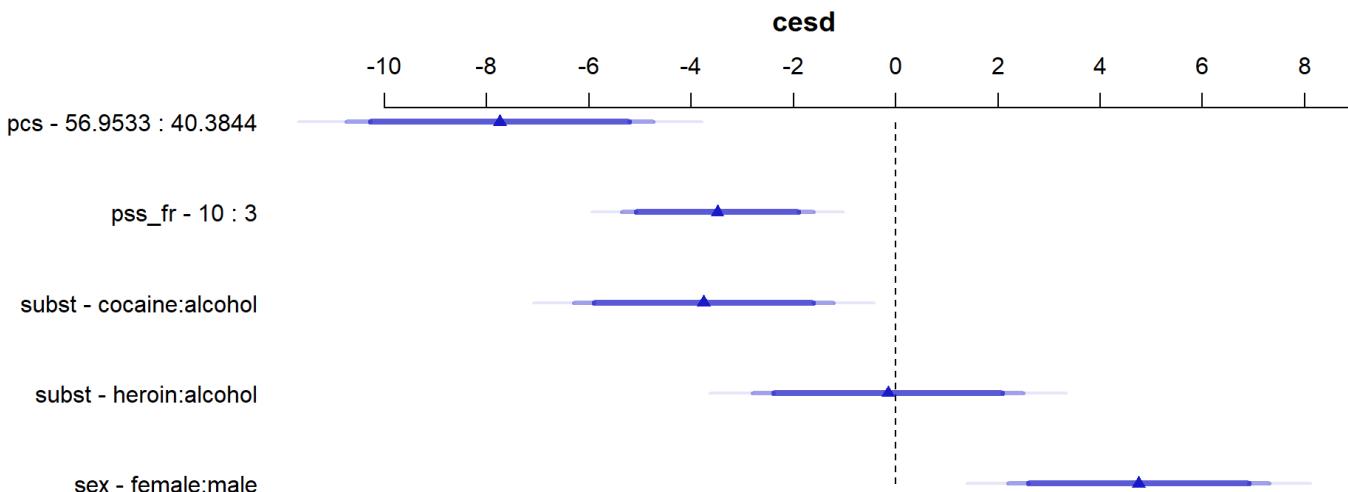
```
Min 1Q Median 3Q Max
```

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

61

Effect Sizes in fitB5

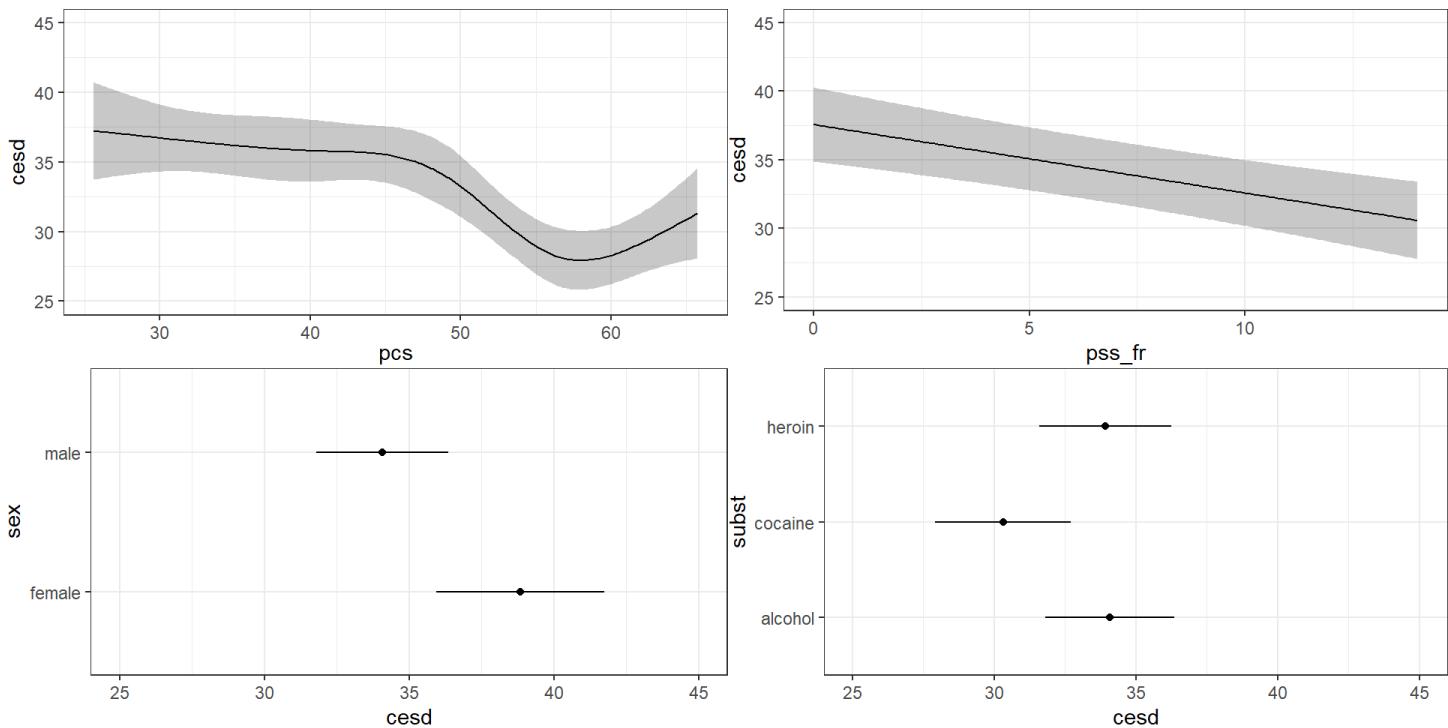
```
1 plot(summary(fitB5))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

62

```
1 ggplot(Predict(fitB5, conf.int = 0.90))
```

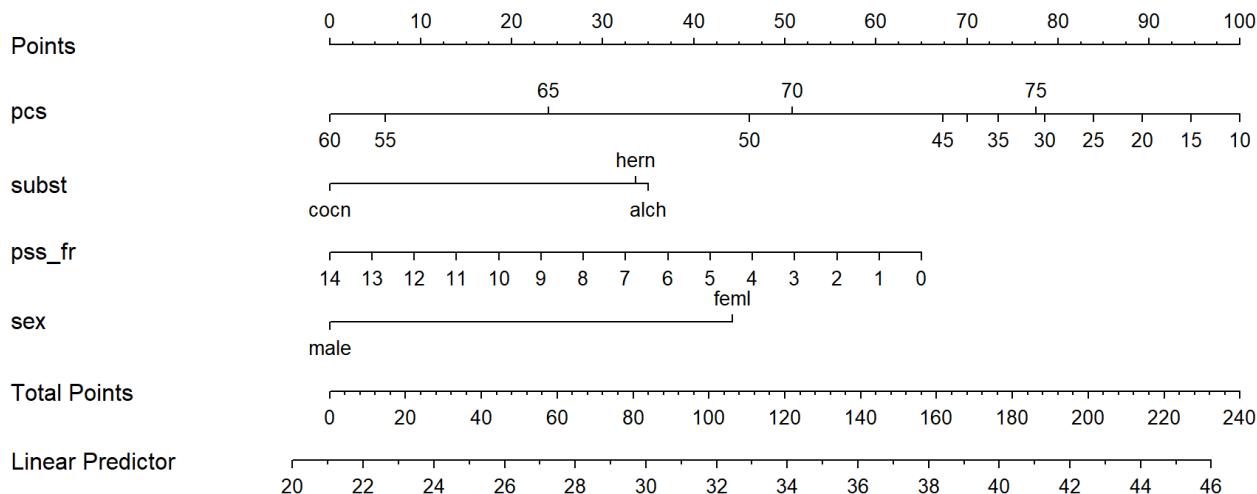


432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

63

A Nomogram for fitB5

```
1 plot(nomogram(fitB5, abbrev = TRUE))
```



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

64

What if you're doing a logistic regression?

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

65

Predicting $\text{Pr}(\text{CESD}>15)$ with a spline

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fitD5 <- lrm(cesd_hi ~ rcs(pcs,5) + subst + pss_fr + sex,
5                 data = help1, x = TRUE, y = TRUE)
6
7 fitD5$coefficients
```

	Intercept	pcs	pcs'	pcs''	pcs'''
6.62332347	6.62332347	-0.05399131	0.03371332	-0.87238631	3.36265431
subst=cocaine	subst=heroin		pss_fr	sex=male	
-1.10386515	-1.10386515	-0.40683747	-0.08278496	-0.22546869	

Contents of fitD5?

1 fitD5

Logistic Regression Model

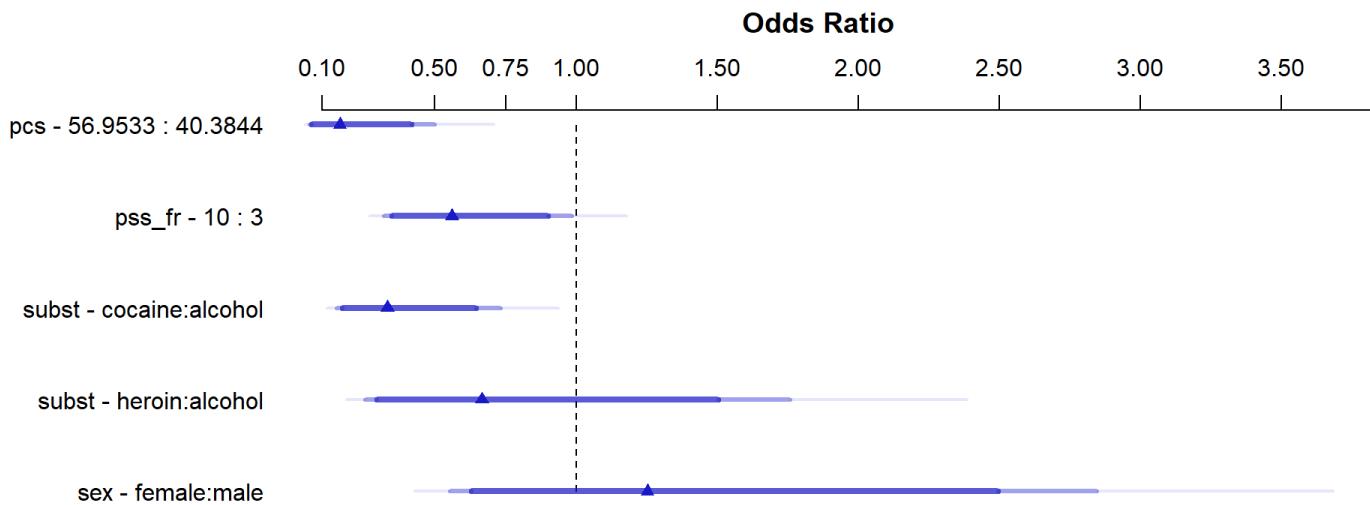
```
lrm(formula = cesd_hi ~ rcs(pcs, 5) + subst + pss_fr + sex, data = help1,  
x = TRUE, y = TRUE)
```

	Model	Likelihood Ratio Test	Discrimination Indexes	Rank Discrim.
Obs	453	LR chi2	42.01	R2 0.184
0	46	d.f.	8	R2(8,453) 0.072
1	407	Pr(> chi2)	<0.0001	R2(8,124) 0.240
max deriv	3e-05			Brier 0.083
				gamma 0.555
				tau-a 0.102
	Coef	S.E.	Wald Z	Pr(> Z)
Intercept	6.6233	4.8970	1.35	0.1762

432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

67

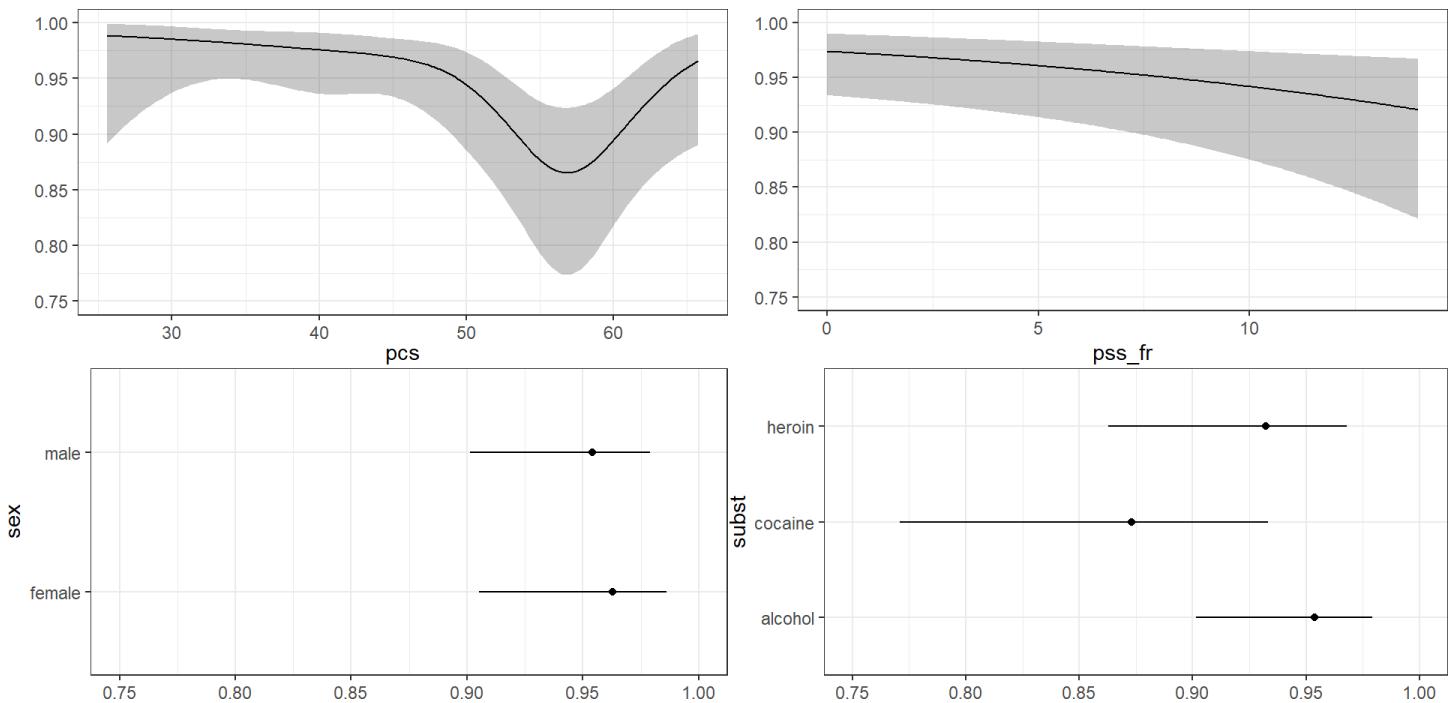
Effect Sizes in fitD5



432 Class 05 | 2026-01-27 | <https://thomaselove.github.io/432-2026/>

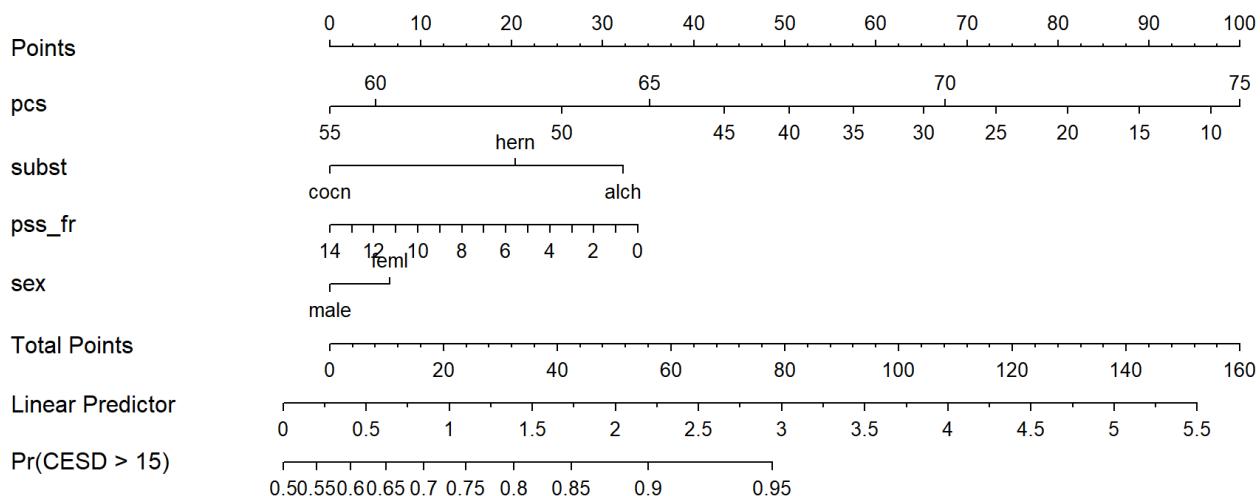
68

```
1 ggplot(Predict(fitD5, conf.int = 0.90, fun = plogis))
```



A Nomogram for fitD5

```
1 plot(nomogram(fitD5, abbrev = TRUE, fun = plogis, funlabel = "Pr(CESD > 15)"))
```



What's next?

- Spearman's ρ^2 plot: exploring non-linearity
 - Spending degrees of freedom on non-linearity wisely
- Model Calibration
- Making Predictions with `ols()` models
- Using both `ols()` and `lm()` as fitters