

432 Class 06

Thomas E. Love, Ph.D.

2026-01-29

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

1

Today's Agenda

- Splines and other non-linear terms
- Spearman's ρ^2 plot: exploring non-linearity
 - Spending degrees of freedom wisely
- Linear Regression (HELP trial again)
 - A complex model with non-linear terms
 - Assessing fit with `ols()` and `lm()`
 - Calibration of the model
 - Prediction Intervals and Confidence Intervals

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

2

Today's R Setup

```
1 knitr::opts_chunk$set(comment = NA)
2
3 library(janitor)
4 library(naniar)
5 library(broom); library(gt); library(patchwork)
6 library(haven)
7 library(rms)           ## auto-Loads Hmisc
8 library(easystats)
9 library(tidyverse)
10
11 theme_set(theme_bw())
```

Types of Splines

- A **linear spline** is a continuous function formed by connecting points (called **knots** of the spline) by line segments.
- A **restricted cubic spline** is a way to build highly complicated curves into a regression equation in a fairly easily structured way.
- A restricted cubic spline is a series of polynomial functions joined together at the knots.
 - Such a spline gives us a way to flexibly account for non-linearity without over-fitting the model.

How complex should our spline be?

Restricted cubic splines can fit many different types of non-linearity. Specifying the number of knots (usually 3, 4 or 5) is all you need to define.

- 3 Knots, 2 degrees of freedom, lets the curve “bend” once.
- 4 Knots, 3 degrees of freedom, lets the curve “bend” twice.
- 5 Knots, 4 degrees of freedom; curve “bends” three times.

A simulated data set

```
1 set.seed(20260129)
2
3 sim_data <- tibble(
4     x = runif(250, min = 10, max = 50),
5     y = 3*(x-30) - 0.3*(x-30)^2 + 0.05*(x-30)^3 +
6         rnorm(250, mean = 500, sd = 70)
7 )
8
9 head(sim_data)

# A tibble: 6 × 2
#>   x     y
#>   <dbl> <dbl>
#> 1 24.4  421.
#> 2 46.9   582.
#> 3 38.6   492.
#> 4 29.1   443.
#> 5 16.4   258.
#> 6 29.2   505.
```

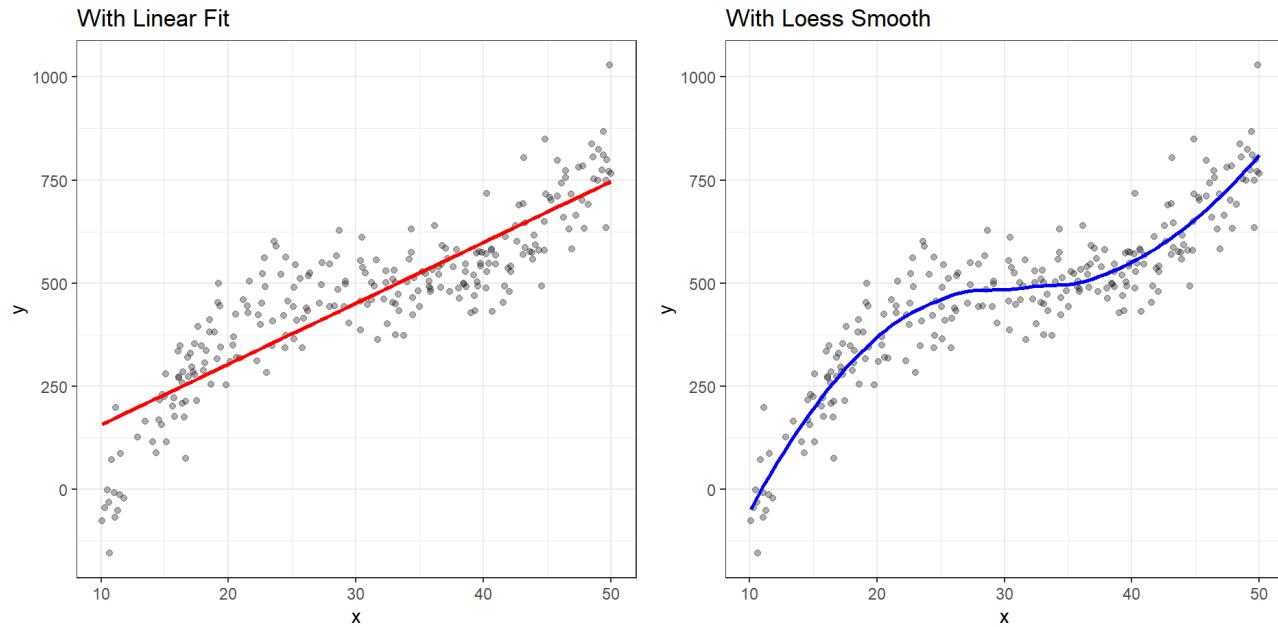
The sim_data, plotted.

```
1 p1 <- ggplot(sim_data, aes(x = x, y = y)) +  
2   geom_point(alpha = 0.3) +  
3   geom_smooth(method = "lm", formula = y ~ x,  
4                 col = "red", se = FALSE) +  
5   labs(title = "With Linear Fit")  
6  
7 p2 <- ggplot(sim_data, aes(x = x, y = y)) +  
8   geom_point(alpha = 0.3) +  
9   geom_smooth(method = "loess", formula = y ~ x,  
10              col = "blue", se = FALSE) +  
11  labs(title = "With Loess Smooth")  
12  
13 p1 + p2
```

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

7

The sim_data, plotted.



432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

8

Fitting Non-Linear Terms with lm

We'll fit:

- a linear model
- two models using orthogonal polynomials (`poly()`), and
- three models using restricted cubic splines (`rcs()`)

```
1 sim_linear <- lm(y ~ x, data = sim_data)
2 sim_poly2 <- lm(y ~ poly(x, 2), data = sim_data)
3 sim_poly3 <- lm(y ~ poly(x, 3), data = sim_data)
4 sim_rcs3 <- lm(y ~ rcs(x, 3), data = sim_data)
5 sim_rcs4 <- lm(y ~ rcs(x, 4), data = sim_data)
6 sim_rcs5 <- lm(y ~ rcs(x, 5), data = sim_data)
```

Degrees of Freedom for each model

- We can check df with `anova(modelname)`

Formula	Model df	Resid. df	# obs.
<code>lm(y ~ x)</code>	1	248	250
<code>lm(y ~ poly(x,2))</code>	2	247	250
<code>lm(y ~ poly(x,3))</code>	3	246	250
<code>lm(y ~ rcs(x,3))</code>	2	247	250
<code>lm(y ~ rcs(x,4))</code>	3	246	250
<code>lm(y ~ rcs(x,5))</code>	4	245	250

augment() for our six models

augment() generates fitted `y` predictions and residuals, which will help us plot the fits for our six models.

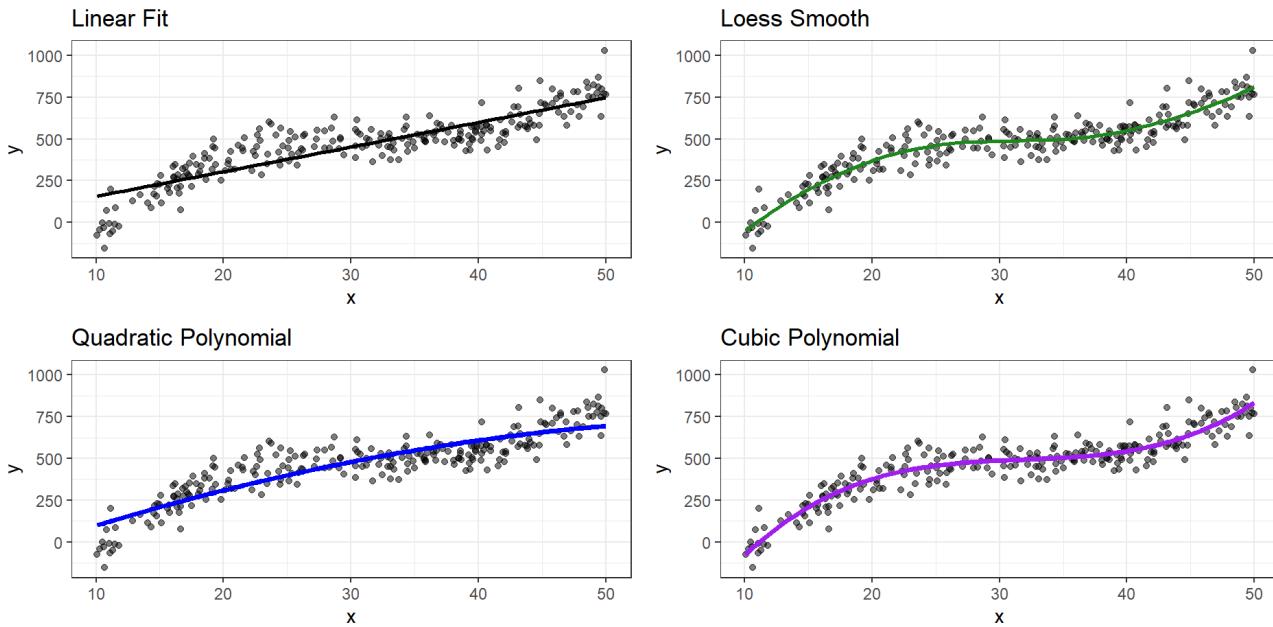
```
1 sim_linear_aug <- augment(sim_linear, sim_data)
2 sim_poly2_aug <- augment(sim_poly2, sim_data)
3 sim_poly3_aug <- augment(sim_poly3, sim_data)
4 sim_rcs3_aug <- augment(sim_rcs3, sim_data)
5 sim_rcs4_aug <- augment(sim_rcs4, sim_data)
6 sim_rcs5_aug <- augment(sim_rcs5, sim_data)
7
8 sim_linear_aug |> slice(1:2) |>
9   gt() |> fmt_number(decimals = 3) |> tab_options(table.font.size = 20)
```

x	y	.fitted	.resid	.hat	.sigma	.cooksdi	.std.resid
24.351	421.499	368.381	53.118	0.005	96.223	0.001	0.554
46.934	582.366	701.362	-118.996	0.012	95.981	0.009	-1.246

Add the Polynomial Fits

```
1 p1 <- ggplot(sim_data, aes(x = x, y = y)) +
2   geom_point(alpha = 0.5) +
3   geom_smooth(method = "lm", formula = y ~ x,
4               col = "black", se = F) +
5   labs(title = "Linear Fit")
6
7 p2 <- ggplot(sim_data, aes(x = x, y = y)) +
8   geom_point(alpha = 0.5) +
9   geom_smooth(method = "loess", formula = y ~ x,
10              col = "forestgreen", se = F) +
11   labs(title = "Loess Smooth")
12
13 p3 <- ggplot(sim_poly2_aug, aes(x = x, y = y)) +
14   geom_point(alpha = 0.5) +
15   geom_line(aes(x = x, y = .fitted),
16             col = "blue", linewidth = 1.25) +
17   labs(title = "Quadratic Polynomial")
18
```

Add the Polynomial Fits



432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

13

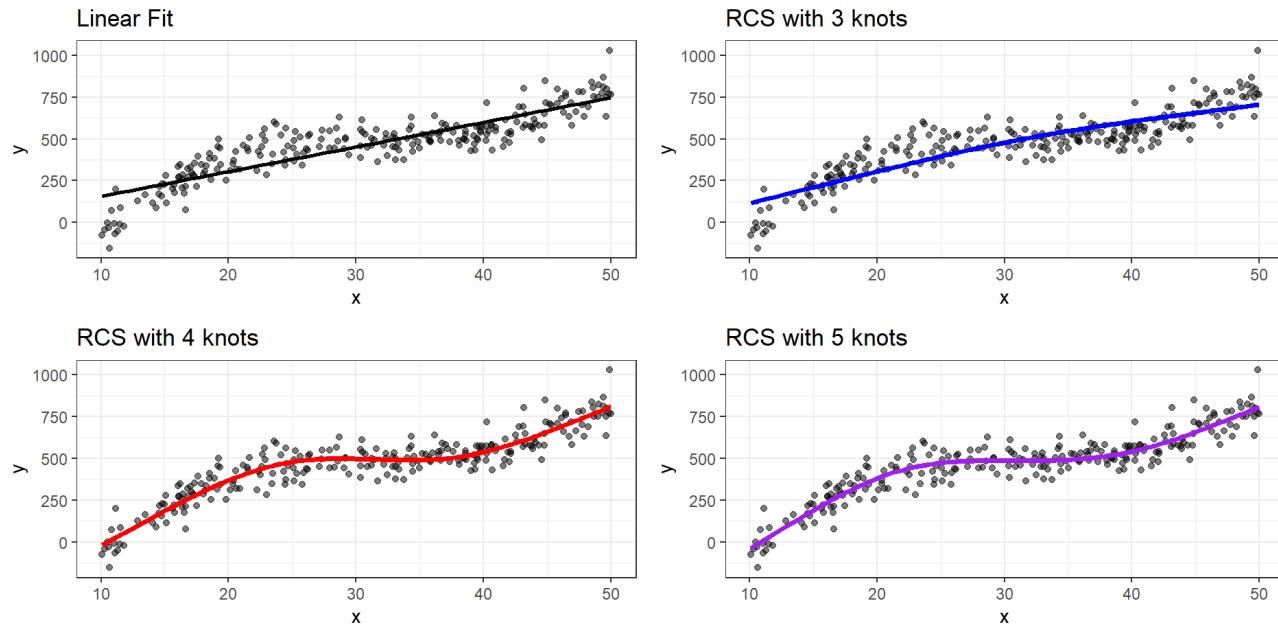
Restricted Cubic Spline Fits

```
1 p0 <- ggplot(sim_data, aes(x = x, y = y)) +  
2   geom_point(alpha = 0.5) +  
3   geom_smooth(method = "lm", formula = y ~ x,  
4               col = "black", se = F) +  
5   labs(title = "Linear Fit")  
6  
7 p3 <- ggplot(sim_rcs3_aug, aes(x = x, y = y)) +  
8   geom_point(alpha = 0.5) +  
9   geom_line(aes(x = x, y = .fitted),  
10            col = "blue", size = 1.25) +  
11   labs(title = "RCS with 3 knots")  
12  
13 p4 <- ggplot(sim_rcs4_aug, aes(x = x, y = y)) +  
14   geom_point(alpha = 0.5) +  
15   geom_line(aes(x = x, y = .fitted),  
16            col = "red", size = 1.25) +  
17   labs(title = "RCS with 4 knots")  
18
```

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

14

Restricted Cubic Spline Fits



Deciding Where to Try Non-Linear Terms

Spending degrees of freedom wisely

- Suppose we have many possible predictors, and minimal theory or subject matter knowledge to guide us.
- We might want our final inferences to be as unbiased as possible. To accomplish this, we have to pay a penalty (in terms of degrees of freedom) for any “peeks” we make at the data in advance of fitting a model.
- So that rules out a lot of decision-making about non-linearity based on looking at the data, if our sample size isn’t incredibly large.

Back to the HELP Trial

Health Evaluation and Linkage to Primary Care (HELP) was a clinical trial of adult inpatients recruited from a detoxification unit.

- We have baseline data for each subject on several variables, including two outcomes:

Variable	Description
----------	-------------

<code>cesd</code>	Center for Epidemiologic Studies-Depression
-------------------	---

<code>cesd_hi</code>	<code>cesd</code> above 15 (indicates high risk)
----------------------	--

help1 data load

```
1 help1 <- tibble(mosaicData::HELPrc) |>
2   select(id, cesd, age, sex, subst = substance, mcs, pcs, pss_fr) |>
3   zap_label() |>
4   mutate(across(where(is.character), as_factor),
5         id = as.character(id),
6         cesd_hi = factor(as.numeric(cesd >= 16)))
7
8 dim(help1); n_miss(help1)
```

```
[1] 453 9
```

```
[1] 0
```

```
1 head(help1, 5)
```

```
# A tibble: 5 × 9
  id    cesd    age sex   subst      mcs    pcs pss_fr cesd_hi
  <chr> <int> <int> <fct> <fct>    <dbl> <dbl>  <int> <fct>
1 1       49     37 male cocaine  25.1    58.4     0 1
2 2       30     37 male alcohol  26.7    36.0     1 1
3 3       39     26 male heroin   6.76   74.8    13 1
4 4       15     39 female heroin  44.0    61.9    11 0
5 5       39     32 male cocaine  21.7    37.3    10 1
```

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

19

The Six Predictors in help1

- Predict cesd using these six predictors...

Variable Description

age subject age (in years)

sex female (n = 107) or male (n = 346)

subst substance abused (alcohol, cocaine, heroin)

mcs SF-36 Mental Component Score

pcs SF-36 Physical Component Score

pss_fr perceived social support by friends

Non-Linear Terms Spends DF

What happens when we add a non-linear term?

- A polynomial of degree D costs D degrees of freedom.
 - So a polynomial of degree 2 (quadratic) costs 2 df, or 1 more than the main effect alone.
- A restricted cubic spline with K knots costs K-1 df.
 - So adding a spline with 4 knots uses 3 df, or 2 more than the main effect alone.
 - We'll only consider splines with 3, 4, or 5 knots.

Non-Linear Terms Spends DF

Adding an interaction (product term) depends on the main effects of the predictors we are interacting

- If the product term's predictors have df_1 and df_2 degrees of freedom, product term adds $df_1 \times df_2$ degrees of freedom.
 - An interaction of a binary and quantitative variable adds $1 \times 1 = 1$ more df to the main effects model.
- When we use a quantitative variable in a spline and interaction, we'll do the interaction on the main effect, not the spline.

A smart first step?

Spearman's is an indicator (not a perfect one) of potential predictive punch, but doesn't give away the game.

- Looking at Spearman's and selecting predictors to include non-linearity for reduces the impact of “looking at the data” which leads to bias in the model.
- Idea: Perhaps we should focus our efforts re: non-linearity on predictors that score better on this measure.

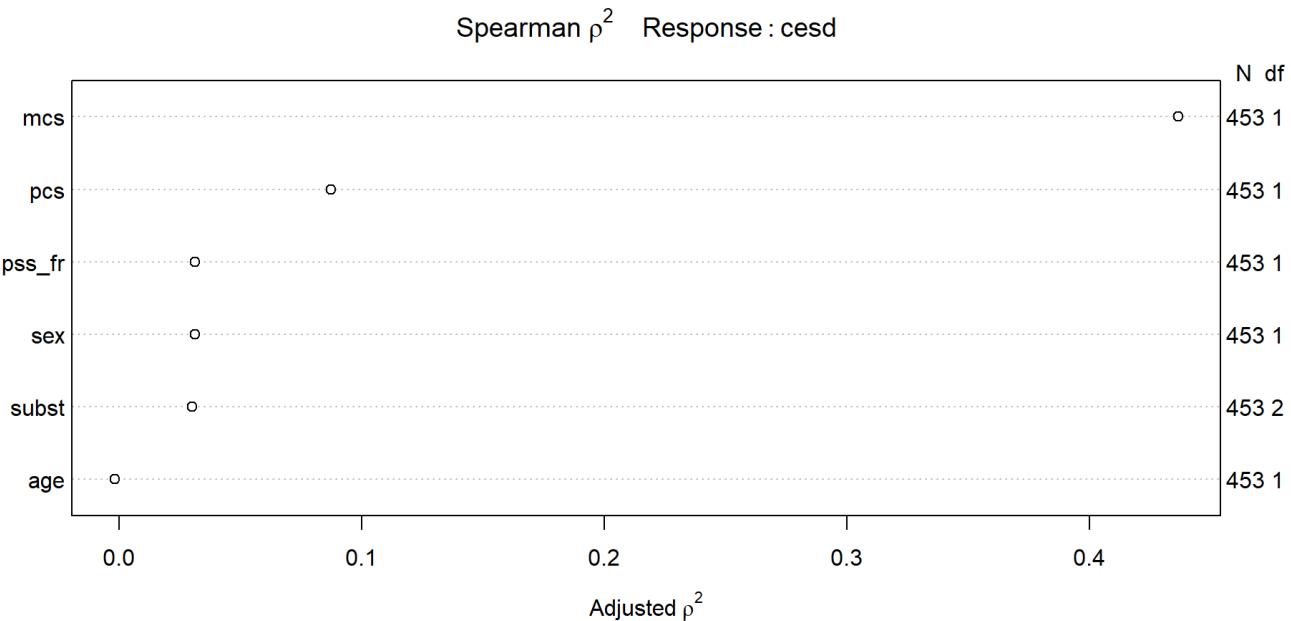
```
1 spear_cesd <- spearman2(cesd ~ mcs + subst + pcs + age + sex + pss_fr,  
2                               data = help1)
```

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

23

Spearman's Plot

```
1 plot(spear_cesd)
```



Conclusions from Spearman Plot

- `mcs` is the most attractive candidate for a non-linear term, as it packs the most potential predictive punch, so if it does turn out to need non-linear terms, our degrees of freedom will be well spent.
 - This **does not** mean that `mcs` actually needs a non-linear term, or will show meaningfully better results if a non-linear term is included. We'd have to fit a model with and without non-linearity in `mcs` to know that.

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

25

Conclusions from Spearman Plot

- `pcs`, also quantitative, has the next most potential predictive punch after `mcs`.
- `pss_fr` and `sex` follow, then `subst` and `age`.

```
1 spear_cesd
```

Spearman rho^2 Response variable:cesd

	rho2	F	df1	df2	P	Adjusted rho2	n
mcs	0.438	350.89	1	451	0.0000	0.436	453
subst	0.034	7.97	2	450	0.0004	0.030	453
pcs	0.089	44.22	1	451	0.0000	0.087	453
age	0.000	0.12	1	451	0.7286	-0.002	453
sex	0.033	15.56	1	451	0.0001	0.031	453
pss_fr	0.033	15.57	1	451	0.0001	0.031	453

A Main Effects Model

Here's a summary of the degrees of freedom for a main effects model without any non-linear terms.

```
1 fit1 <- lm(cesd ~ mcs + subst + pcs + age + sex + pss_fr, data = help1)
2
3 glance(fit1) |> select(df, df.residual, nobs) |>
4   gt() |> tab_options(table.font.size = 20) |>
5   opt_stylize(style = 3, color = "cyan")
```

df	df.residual	nobs
7	445	453

We started with 453 observations (452 df) and fitting `fit1` leaves 445 residual df, so `fit1` uses 7 degrees of freedom.

Grim Reality

One popular standard for linear regression requires at least 25 observations *per regression coefficient that you will estimate*¹.

- With 453 observations (452 df) in the HELP trial, we should be thinking about models with modest numbers of regression inputs, since 25 is really a bare minimum.
- We've already committed to 7 such coefficients (intercept + our six predictors.)

Sample Size (spending df)

- Non-linear terms (polynomials, splines, product terms) just add to the problem, as they need additional degrees of freedom (parameters) to be estimated.
- We'll also use more df every time if we consider re-fitting after variable selection.

So we might choose to include non-linear terms in just two or three variables with this modest sample size ($n = 453$).

- But I'll ignore all of that (for now) and propose a complex `fit2` model ...

Proposed New Model `fit2`

Fit a model to predict `cesd` using:

- a 5-knot spline on `mcs`
- a 3-knot spline on `pcs`
- a polynomial of degree 2 on `pss_fr`
- a linear term on `age`
- an interaction of `sex` with the main effect of `mcs` (restricting our model so that terms that are non-linear in both `sex` and `mcs` are excluded), and
- a main effect of `subst`

Standing Break

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

31

Our new model `fit2`

Definitely more than we can reasonably do with 453 observations, but let's see how it looks.

```
1 dd <- datadist(help1)
2 options(datadist = "dd")
3
4 fit2 <- ols(cesd ~ rcs(mcs, 5) + rcs(pcs, 3) + sex + mcs %ia% sex +
5             pol(pss_fr,2) + age + subst,
6             data = help1, x = TRUE, y = TRUE)
```

- `%ia%` tells R to fit an interaction term with `sex` and the main effect of `mcs`.
 - We have to include `sex` as a main effect for the interaction term (`%ia%`) to work. We already have the main effect of `mcs` in as part of the spline.

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

32

Can we fit2 with lm()?

Yes. Note poly() in our lm() fit, rather than pol().

```
1 fit2_lm <- lm(cesd ~ rcs(mcs, 5) + rcs(pcs, 3) + sex + mcs %ia% sex +
2                         poly(pss_fr, 2) + age + subst, data = help1)
3
4 glance(fit2_lm) |> select(df, df.residual, nobs) |>
5   gt() |> tab_options(table.font.size = 20) |>
6   opt_stylize(style = 3, color = "cyan")
```

df	df.residual	nobs
13	439	453

- So fit2_lm uses an additional 6 degrees of freedom beyond the 7 in fit1.

Our fitted model fit2 (from ols())

```
1 fit2
```

Linear Regression Model

```
ols(formula = cesd ~ rcs(mcs, 5) + rcs(pcs, 3) + sex + mcs %ia%
  sex + pol(pss_fr, 2) + age + subst, data = help1, x = TRUE,
  y = TRUE)
```

	Model Likelihood	Discrimination	
	Ratio Test	Indexes	
Obs	453	LR chi2 353.70	R2 0.542
sigma8.5942	d.f. 13	R2 adj 0.528	
d.f.	439	Pr(> chi2) 0.0000	g 10.483

Residuals

ANOVA for fit2

This ANOVA testing is sequential, other than the TOTALS.

```
1 anova(fit2)
```

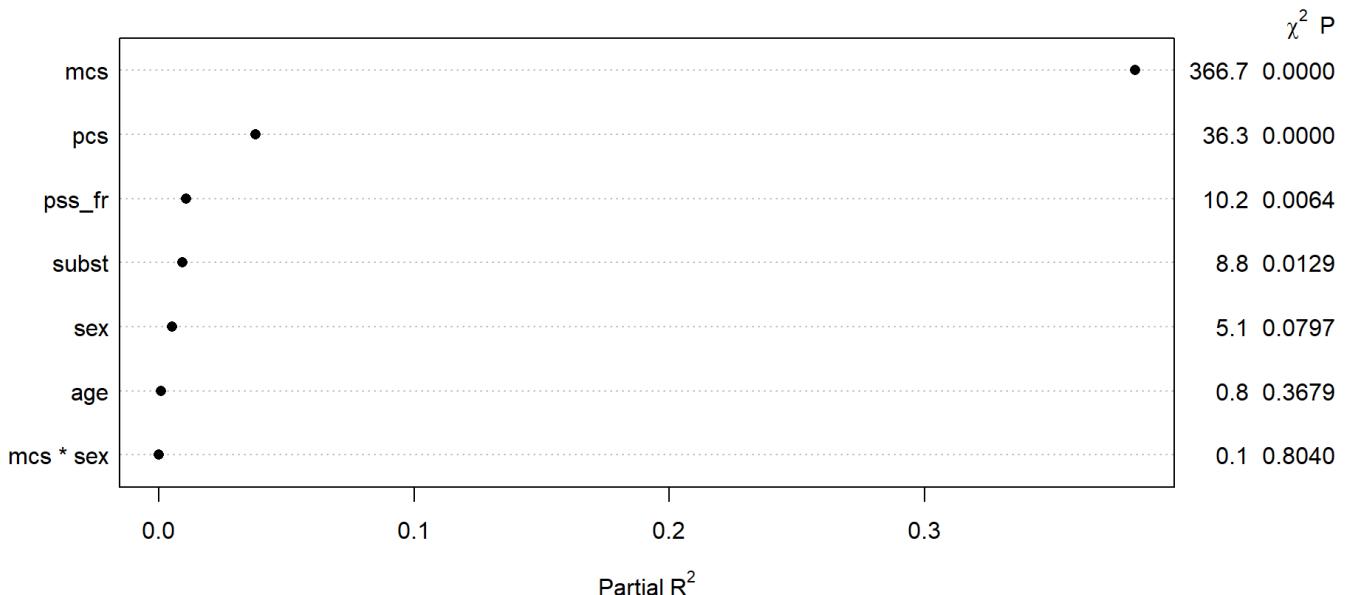
Analysis of Variance		Response: cesd			
Factor		d.f.	Partial SS	MS	F P
mcs (Factor+Higher Order Factors)	All Interactions	5	26857.364671	5371.472934	72.21 <.0001
Nonlinear		1	2.026255	2.026255	0.03 0.8690
pcs		3	293.502251	97.834084	1.32 0.2688
Nonlinear		2	2548.388579	1274.194290	17.13 <.0001
sex (Factor+Higher Order Factors)	All Interactions	2	451.578352	225.789176	3.04 0.0491
mcs * sex (Factor+Higher Order Factors)		1	2.026255	2.026255	0.03 0.8690
pss_fr		1	448.812293	448.812293	6.03 0.0144
age		1	49.758786	49.758786	0.67 0.4139
subst		2	611.625952	305.812976	4.11 0.0170

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

35

Plotting ANOVA results for fit2

```
1 plot(anova(fit2), what = "partial R2", sort = "ascending")
```



432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

36

Validation of Summary Statistics

```
1 set.seed(432); validate(fit2, method = "boot", B = 300)
```

	index.orig	training	test	optimism	index.corrected	Lower	Upper
R-square	0.5420	0.5570	0.5259	0.0311	0.5108	0.4445	0.5735
MSE	71.5769	68.8206	74.0829	-5.2623	76.8393	68.1247	85.9959
g	10.4826	10.5784	10.3304	0.2480	10.2346	9.1147	11.2645
Intercept	0.0000	0.0000	0.7935	-0.7935	0.7935	-3.2042	4.4715
Slope	1.0000	1.0000	0.9753	0.0247	0.9753	0.8738	1.0878
n							
R-square	300						
MSE	300						
g	300						
Intercept	300						
Slope	300						

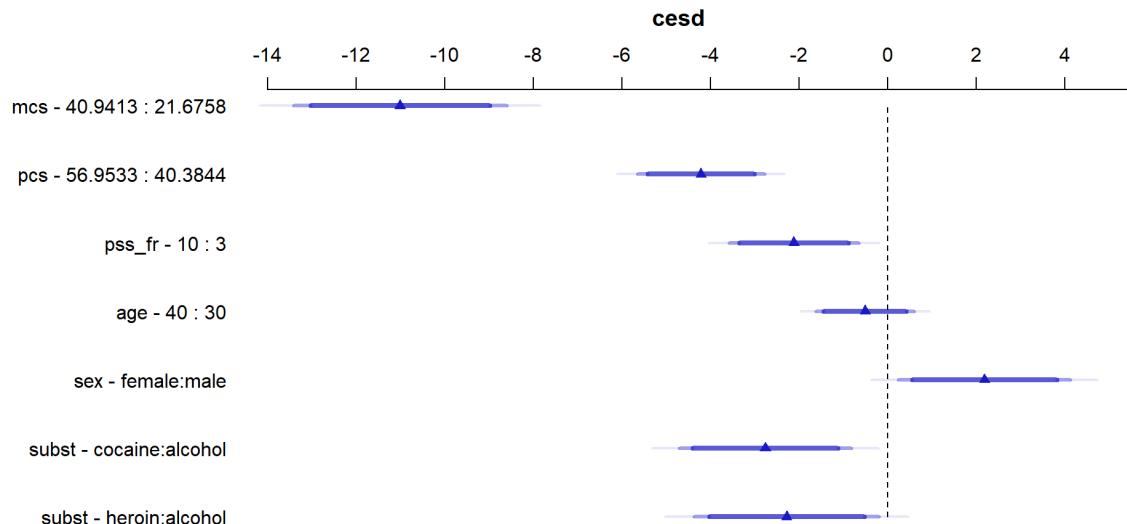
- I'm making a blanket recommendation that you run 300 bootstrap validations unless (in a Lab or something) I've told you specifically to do something else.

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

37

summary results for fit2

```
1 plot(summary(fit2))
```



Adjusted to:mcs=28.60242 sex=males

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

38

summary results for fit2

```
1 summary(fit2)
```

Effects	Response : cesd							
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper
mcs	21.676	40.941	19.266	-11.01300	1.22920	-13.42900	-8.59710	
pcs	40.384	56.953	16.569	-4.21690	0.73316	-5.65780	-2.77590	
pss_fr	3.000	10.000	7.000	-2.12120	0.74667	-3.58870	-0.65369	
age	30.000	40.000	10.000	-0.51164	0.56762	-1.62720	0.60394	
sex - female:male	2.000	1.000	NA	2.18360	0.99288	0.23218	4.13500	
subst - cocaine:alcohol	1.000	2.000	NA	-2.76380	0.99343	-4.71630	-0.81134	
subst - heroin:alcohol	1.000	3.000	NA	-2.28280	1.06530	-4.37640	-0.18915	

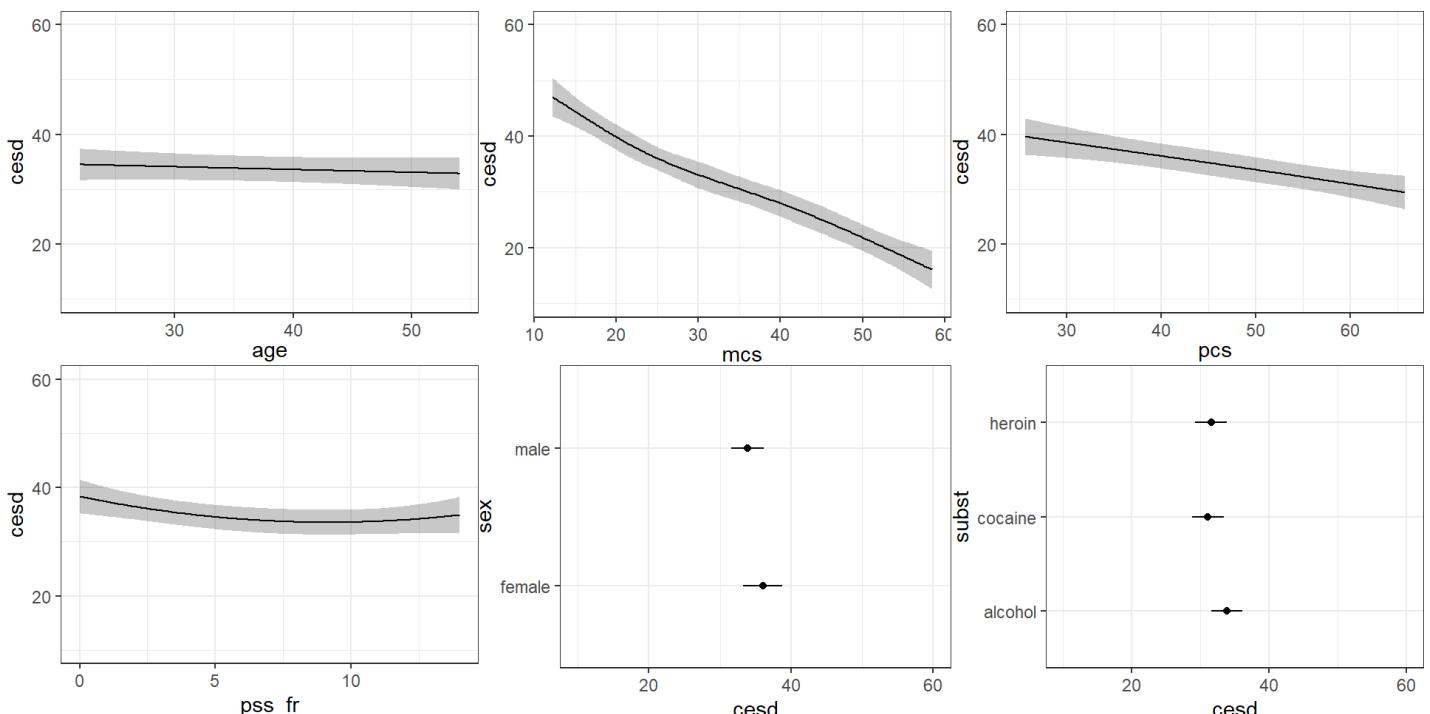
Adjusted to: mcs=28.60242 sex=male

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

39

Impact of non-linearity?

```
1 ggplot(Predict(fit2))
```

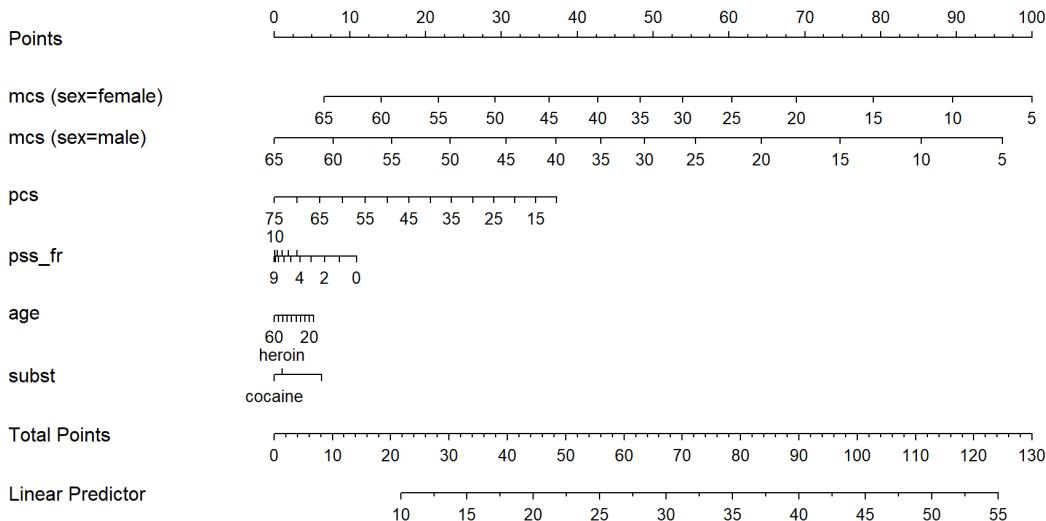


432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

40

Nomogram for fit2

```
1 plot(nomogram(fit2))
```



How to use the nomogram

1. Find the value of each predictor on its provided line, and identify the “points” for that predictor by drawing a vertical line up to the “Points”.
2. Then sum up the points over all predictors to obtain “Total Points”.
3. Draw a vertical line from “Total Points” to “Linear Predictor” to obtain predicted `cesd`.

The nomogram shows modeled effects and their impact on the predicted outcome.

Making Predictions

Suppose we want to use our model `fit2` to make a prediction for `cesd` for a new subject, named Grace, who has the following characteristics...

- sex = female, mcs = 40, pcs = 50
- pss_fr = 7, age = 45, subst = “cocaine”

We can build point and interval estimates for predicted `cesd` from `fit2` as follows...

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

43

Predictions for an Individual

Suppose we have a new **individual subject** named Grace.

```
1 grace <- tibble(sex = "female", mcs = 40, pcs = 50,
2                   pss_fr = 7, age = 45, subst = "cocaine")
3
4 predict(fit2, newdata = grace, conf.int = 0.95, conf.type = "individual") |>
5   as_vector()
```

linear.predictors.1	lower.1	upper.1
26.808537	9.595825	44.021249

Our predicted `cesd` for Grace is 26.81, with 95% **prediction interval** (9.60, 44.02).

Predictions for a Long-Run Mean

Predict mean `cesd` of a set of subjects with Grace's predictor values, along with a **confidence interval**.

```
1 predict(fit2, newdata = grace, conf.int = 0.95, conf.type = "mean") |>  
2 as_vector()
```

linear.predictors.1	lower.1	upper.1
26.80854	23.49523	30.12185

- Confidence interval (23.50, 30.12) is much narrower than prediction interval (9.60, 44.02).

Assessing the Calibration of `fit2`

We would like our model to be well-calibrated, in the following sense...

- Suppose our model assigns a predicted outcome of 6 to several subjects.
- If the model is well-calibrated, this means we expect the mean of those subjects' actual outcomes to be very close to 6.
- We'd like to look at the relationship between the observed `cesd` outcome and our predicted `cesd` from the model.

Building a Calibration Plot

- The calibration plot we'll create provides two estimates (with and without bias-correction) of the predicted vs. observed values of our outcome, and compares these to the ideal scenario (predicted = observed).
- The plot uses resampling validation to produce bias-corrected estimates and uses lowess smooths to connect across predicted values.
- Calibration plots require `x = TRUE, y = TRUE` in a fit with `ols()`.

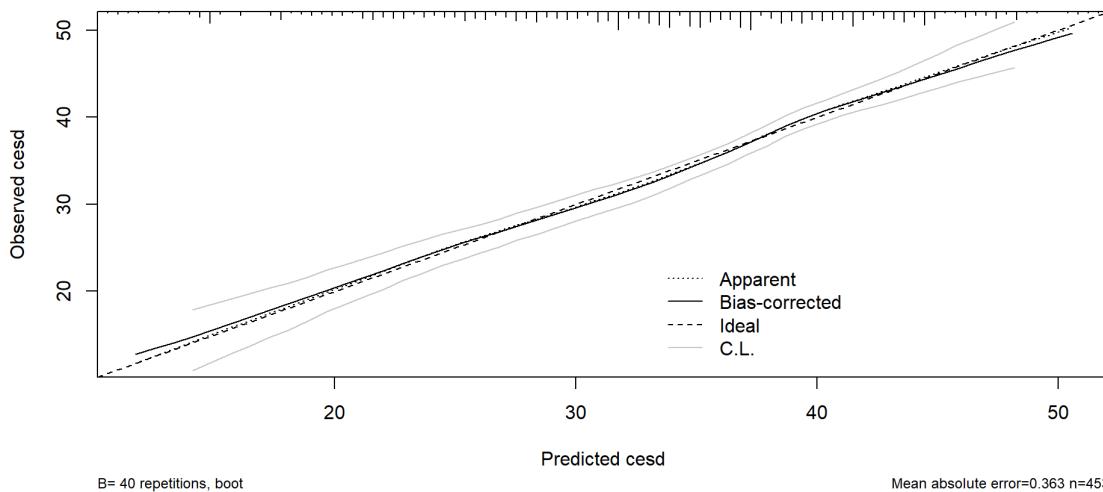
432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

47

Checking the model's calibration

```
1 set.seed(432); plot(calibrate(fit2))
```

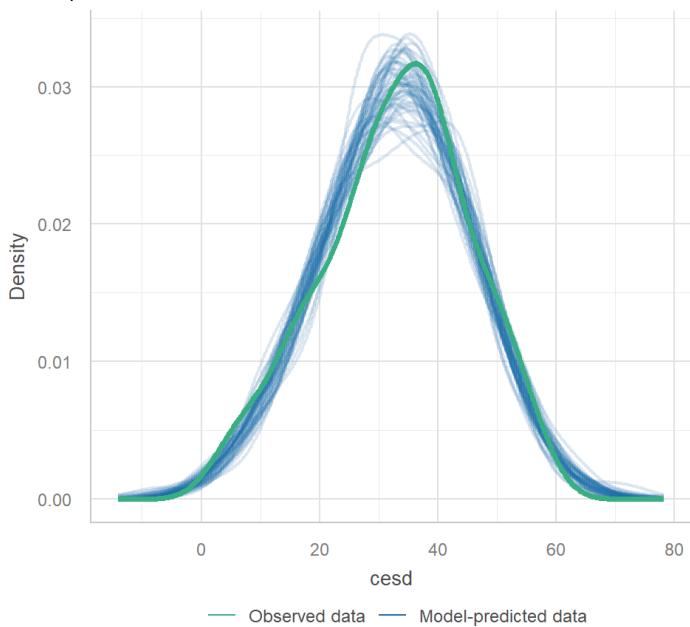
n=453 Mean absolute error=0.363 Mean squared error=0.17069
0.9 Quantile of absolute error=0.627



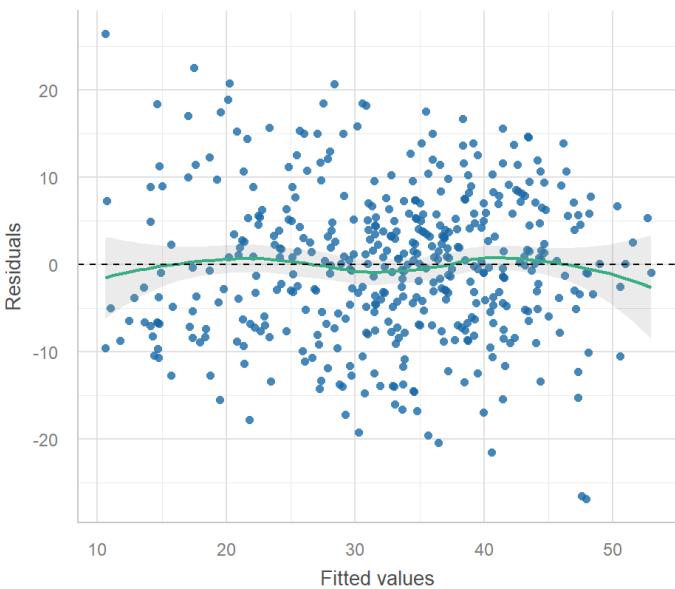
Checking the Model (first 2 plots)

```
1 check_model(fit2_lm, check = c("pp_check", "linearity"))
```

Posterior Predictive Check
Model-predicted lines should resemble observed data line



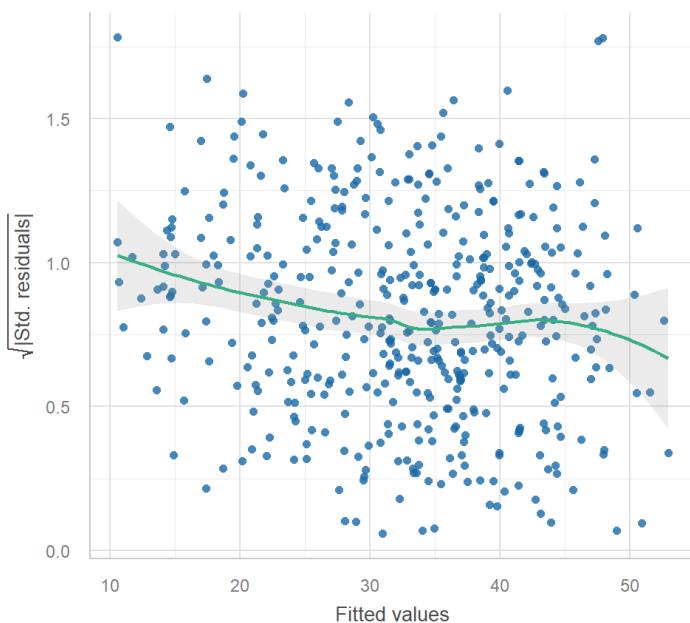
Linearity
Reference line should be flat and horizontal



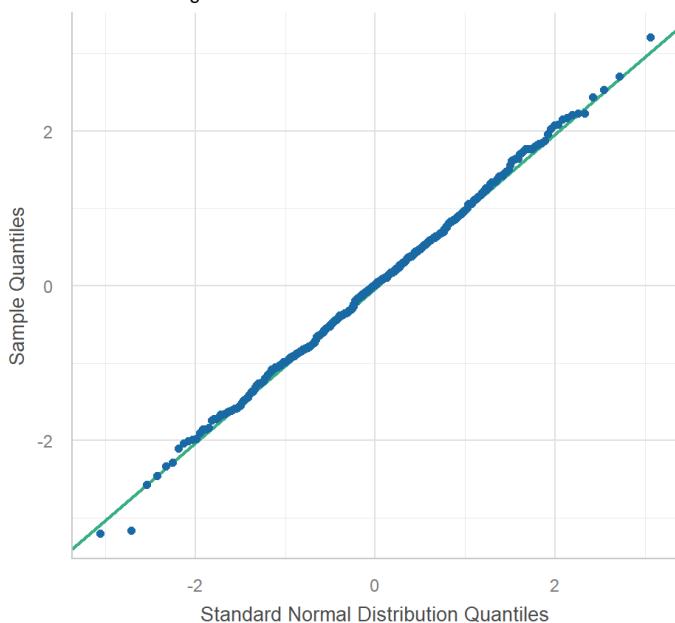
Checking the Model (plots 3-4)

```
1 check_model(fit2_lm, detrend = FALSE, check = c("homogeneity", "qq"))
```

Homogeneity of Variance
Reference line should be flat and horizontal

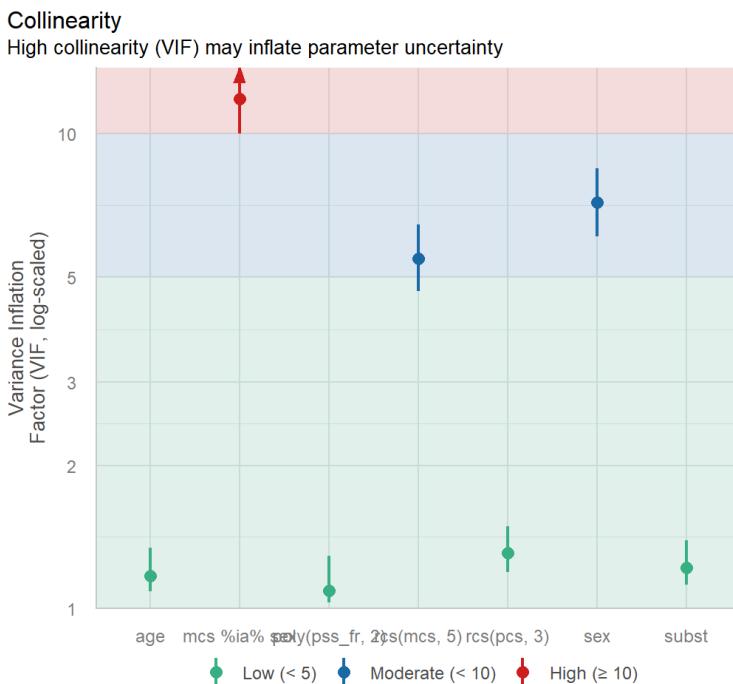


Normality of Residuals
Dots should fall along the line



Checking the model (plot 5)

```
1 check_model(fit2_lm, check = c("vif"))
```



432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

51

Checking Collinearity

```
1 check_collinearity(fit2_lm)
```

```
# Check for Multicollinearity
```

Low Correlation

	Term	VIF	VIF	95% CI	adj. VIF	Tolerance	Tolerance	95% CI
	rcs(pcs, 3)	1.31	[1.19,	1.49]	1.14	0.76	[0.67,	0.84]
	poly(pss_fr, 2)	1.09	[1.03,	1.29]	1.04	0.92	[0.78,	0.97]
	age	1.17	[1.09,	1.34]	1.08	0.85	[0.74,	0.92]
	subst	1.22	[1.12,	1.39]	1.05	0.82	[0.72,	0.89]

Moderate Correlation

	Term	VIF	VIF	95% CI	adj. VIF	Tolerance	Tolerance	95% CI
	rcs(mcs, 5)	5.46	[4.66,	6.43]	2.34	0.18	[0.16,	0.21]

Variance Inflation Factors

The collinearity plot is a bit hard to see with all of these terms, so we can just look at the variance inflation factors:

```
1 rms::vif(fit2)
```

	mcs	mcs'	mcs''	mcs'''	pcs
	53.711079	4838.370091	12475.902431	2489.147506	5.521090
	pcs'	sex=male	mcs * sex=male	pss_fr	pss_fr^2
	5.365910	7.163012	11.848760	15.657046	15.885078
	age	subst=cocaine	subst=heroin		
	1.172137	1.349517	1.383641		

```
1 car::vif(fit2_lm)
```

	GVIF	Df	GVIF^(1/(2*Df))
rcs(mcs, 5)	5.461428	4	1.236412
rcs(pcs, 3)	1.308116	2	1.069453
sex	7.163012	1	2.676380
mcs %ia% sex	11.848760	1	3.442203
poly(pss_fr, 2)	1.091682	2	1.022172
age	1.172137	1	1.082653
subst	1.217443	2	1.050418

432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

53

Tests instead of plots?

- Never, ever, but ...

```
1 check_heteroscedasticity(fit2_lm)
```

Warning: Heteroscedasticity (non-constant error variance) detected (p = 0.048).

```
1 check_normality(fit2_lm)
```

OK: residuals appear as normally distributed (p = 0.986).

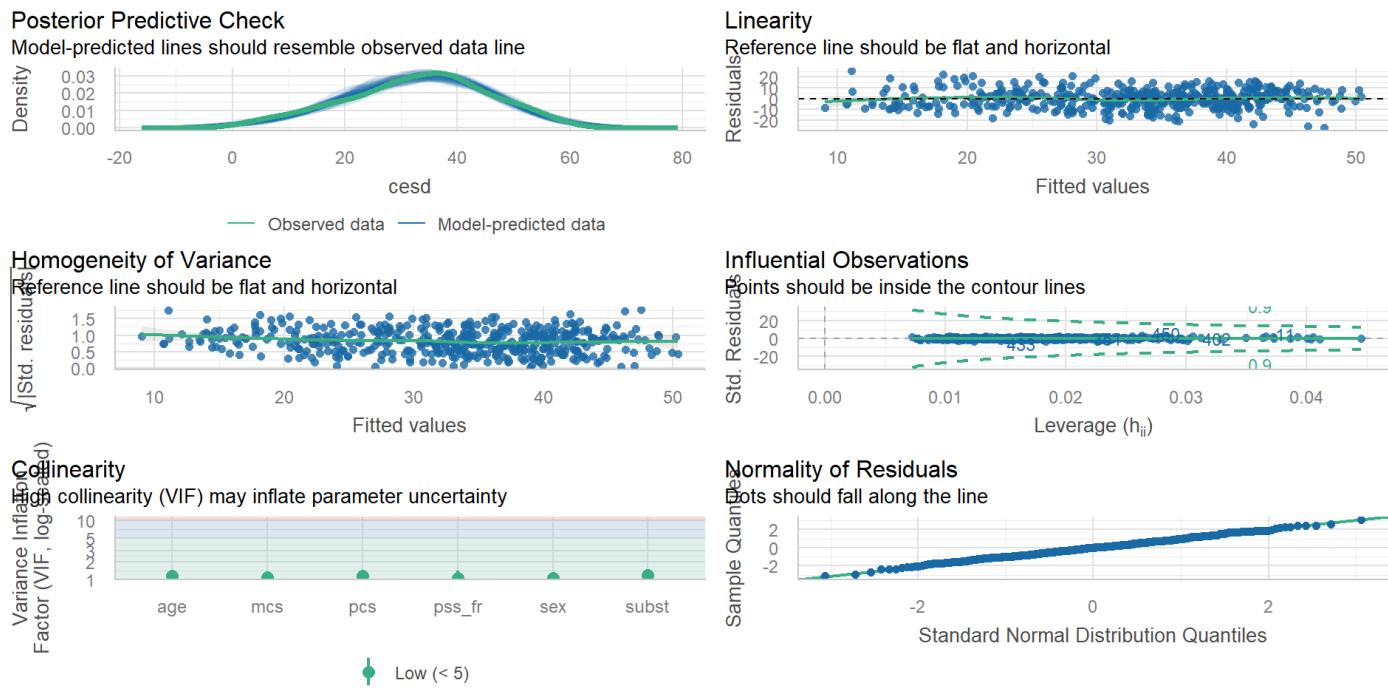
```
1 check_outliers(fit2_lm)
```

OK: No outliers detected.

- Based on the following method and threshold: cook (0.9).
- For variable: (Whole model)

Checking model `fit1`?

```
1 check_model(fit1, detrend = FALSE)
```



432 Class 06 | 2026-01-29 | <https://thomaselove.github.io/432-2026/>

55

Using both `lm()` and `ols()`

- We can and will regularly use both `lm` and `ols` to fit a model like `fit2`.

To delve into the details of how well this complex model works, and to help plot what is actually being fit, we'll want to fit the model using `ols()`.

- In Project A, we expect some results that are most easily obtained using `lm()` and others that are most easily obtained using `ols()`.

What's Coming Up?

- Focus on logistic regression with a new data set
 - Thinking about various pseudo- approaches
 - Developing an optimal cutpoint for a confusion matrix
 - Brier scores and other measures of calibration in logistic regression
 - Checking assumptions in logistic regression
 - Just about everything we might want to do...